

Midterm Review

Summary of Course Material

Basics of neural networks:


- Loss function & Regularization
- Optimization
- Activation Functions

How we build complex network models


- Convolutional Layers

KNN


- How does it work? Train? Test?
- What positive / negative effect would using larger / smaller k value have?
- Distance function? L1 vs. L2?



$K = 1$




$K = 3$



$K = 5$


Linear Classifier



$$f(x, W) = Wx + b$$


Algebraic Viewpoint

$$f(x, W) = Wx$$




Visual Viewpoint

One template
per class



Geometric Viewpoint

Hyperplanes
cutting up space



Suppose: 3 training examples, 3 classes.
With some W the scores $f(x, W) = Wx$ are:



| | | | |
|------|------------|------------|-------------|
| cat | 3.2 | 1.3 | 2.2 |
| car | 5.1 | 4.9 | 2.5 |
| frog | -1.7 | 2.0 | -3.1 |

A **loss function** tells how good our current classifier is

Given a dataset of examples

$$\{(x_i, y_i)\}_{i=1}^N$$

Where x_i is image and
 y_i is (integer) label

Loss over the dataset is a
average of loss over examples:


$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

Multiclass SVM Loss

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases}$$
$$= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$



$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$



cat

3.2

car

5.1

frog

-1.7

Losses:

2.9

$$= \max(0, 5.1 - 3.2 + 1) + \max(0, -1.7 - 3.2 + 1)$$

$$= \max(0, 2.9) + \max(0, -3.9)$$

$$= 2.9 + 0$$

$$= 2.9$$

Softmax Classifier



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{Softmax Function}$$

Probabilities
must be ≥ 0

Probabilities
must sum to 1

$$L_i = -\log P(Y = y_i|X = x_i)$$

cat
car
frog

3.2
5.1
-1.7

Unnormalized
log-probabilities / logits

exp

24.5
164.0
0.18

unnormalized
probabilities

normalize

0.13
0.87
0.00

probabilities

$$\rightarrow L_i = -\log(0.13) = 2.04$$

Loss Intuition:

- Maximizing Likelihood
- Comparing Probabilities

Regularization

$$f(x, W) = Wx$$

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1)$$

E.g. Suppose that we found a W such that $L = 0$.
Is this W unique?

No! $2W$ also has $L = 0$!

How do we choose between W and $2W$?

Regularization

Regularization

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

Data loss: Model predictions should match training data

Regularization: Prevent the model from doing *too* well on training data

Why regularize?

- Express preferences over weights
- Make the model *simpler* to avoid overfitting

Optimization Motivation

- We have some dataset of (x,y)
- We have a **score function**:
- We have a **loss function**:

$$s = f(x; W) = Wx$$

$$L = \frac{1}{N} \sum_{i=1}^N L_i + R(W)$$

How do we find the best W ?

Gradient Descent

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(x_i, y_i, W) + \lambda R(W)$$

$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

Full sum expensive
when N is large!

Approximate sum
using a **minibatch** of
examples
32 / 64 / 128 common

```
# Vanilla Minibatch Gradient Descent
```


```
while True:
```

```
    data_batch = sample_training_data(data, 256) # sample 256 examples
```

```
    weights_grad = evaluate_gradient(loss_fun, data_batch, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```


Optimization, Point 1: Calculating Gradients for Updates




Computational Graphs + Backpropagation (Chain Rule)

Optimization, Point 2: Things to take care!



Local Minima



Saddle points



Poor Conditioning




Other Algos: SGD+momentum, AdaGrad, RMSProp, Adam

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Convolution Layer




For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:
activation maps



We stack these up to get a “new image” of size 28x28x6!

Convolution Layer

In contrast to fully connected layer,
Each term in output is dependent on spatially local 'subregions' of input




$$out_i = \sum_{j=1}^{H \times W \times C} w_{ij} \cdot in_j + b_i$$

Convolution Layer

In contrast to fully connected layer,

Each term in output is dependent on spatially local 'subregions' of input



Question: connection between an FC layer and a convolutional layer?


$$out_i = \sum_{j=1}^{H \times W \times C} w_{ij} \cdot in_j + b_i$$

$$out_i = \sum_{j=1}^{HH \times WW \times C} w_{ij} \cdot in(patch_i)_j + b_i$$

Convolution Layer

In contrast to fully connected layer,

Each term in output is dependent on spatially local 'subregions' of input



Question: connection between an FC layer and a convolutional layer?

Answer: FC looks like convolution layer with filter size $H \times W$

$$out_i = \sum_{j=1}^{H \times W \times C} w_{ij} \cdot in_j + b_i$$

$$out_i = \sum_{j=1}^{HH \times WW \times C} w_{ij} \cdot in(patch_i)_j + b_i$$

Convolution Layer

For kernel width \mathbf{k} and stride \mathbf{s} ,
Input width \mathbf{w}_{in} and total padding \mathbf{w}_{pad} ,
Output width \mathbf{w}_{out} is

$$w_{out} = \frac{1}{s}(w_{in} + w_{pad} - k) + 1$$