Multi-Modal AI

Chuang Gan

Learning from images and video frames

BSD (2001)

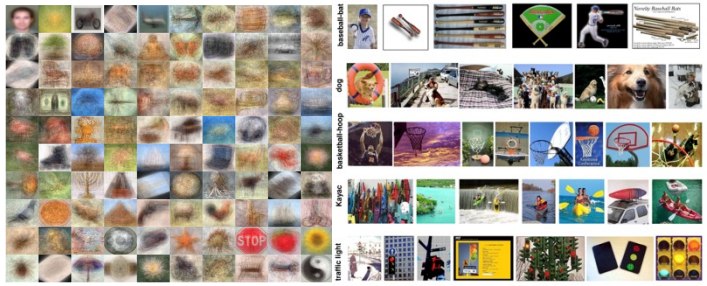

Caltech 101 (2004), Caltech 256 (2006)

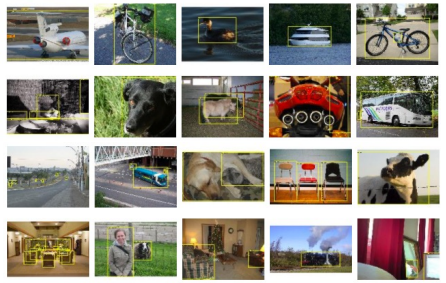

PASCAL (2007-12)

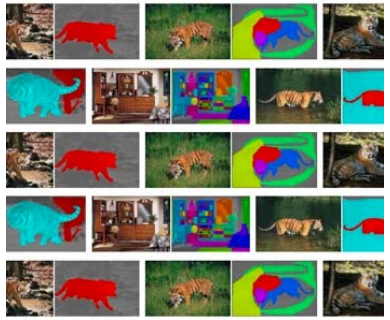

LabelMe (2007)

ImageNet (2009)

SUN (2010)

UCF-101 (2012)

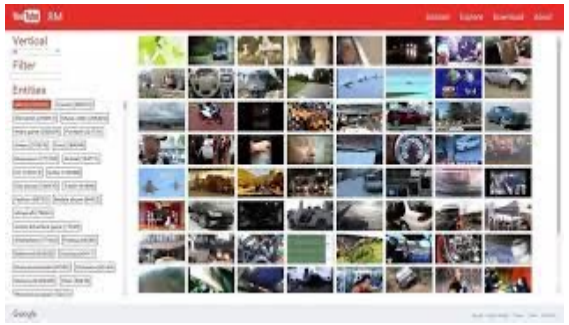

Youtube-8M (2017)

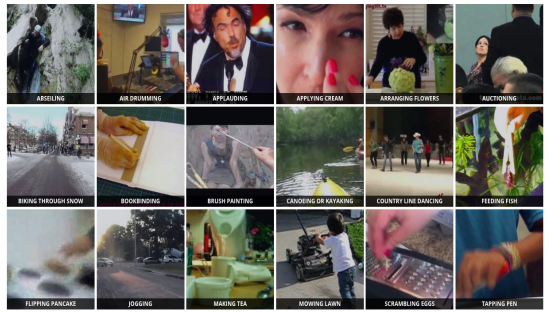

Kinetics (2017)

What can sound give us?

physical interactions

speech

sound of distant object

Can machines connect sight with sound for rich perception?

Task: Visual Sound Separation

Given a music performance video…

Mixed sound

Task: Visual Sound Separation

…we aim to separate two sounds played by different instruments.

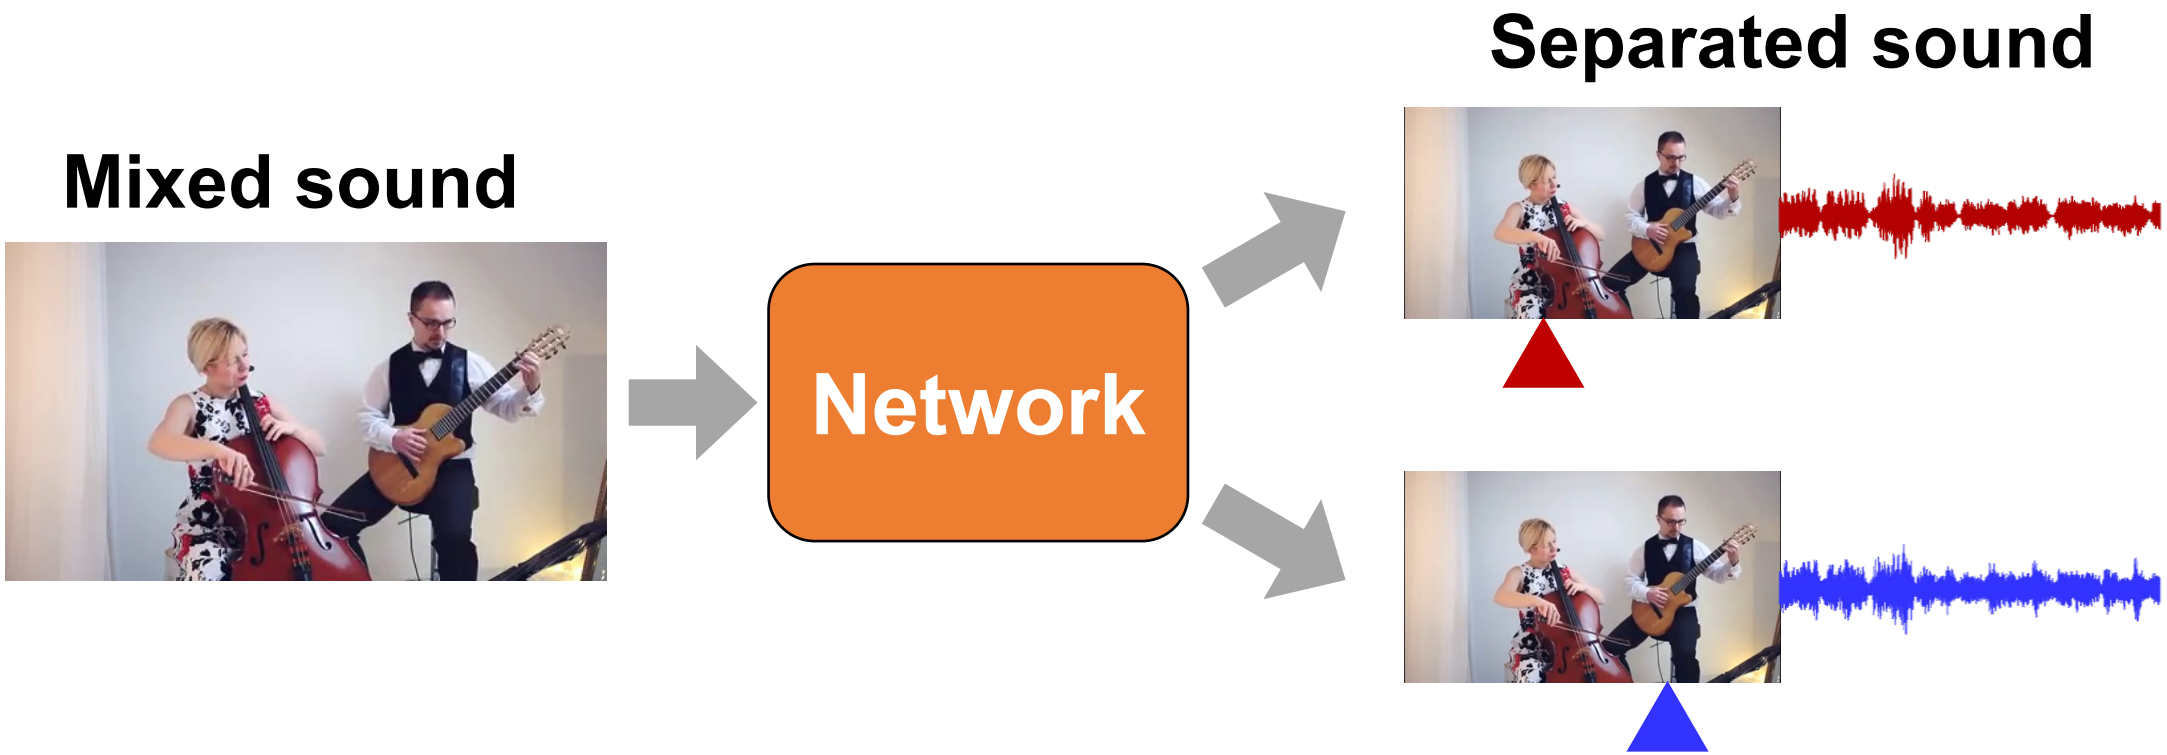

Source Separation: Traditional Approach

- To separate one component out of K:
  - 1. Estimate the mask for the target component

$$\mathbf{M}_s = \frac{\mathbf{w}_1 \, \mathbf{h}_1^{\mathrm{T}}}{\sum_{i=1}^{K} \mathbf{w}_i \, \mathbf{h}_i^{\mathrm{T}}}$$

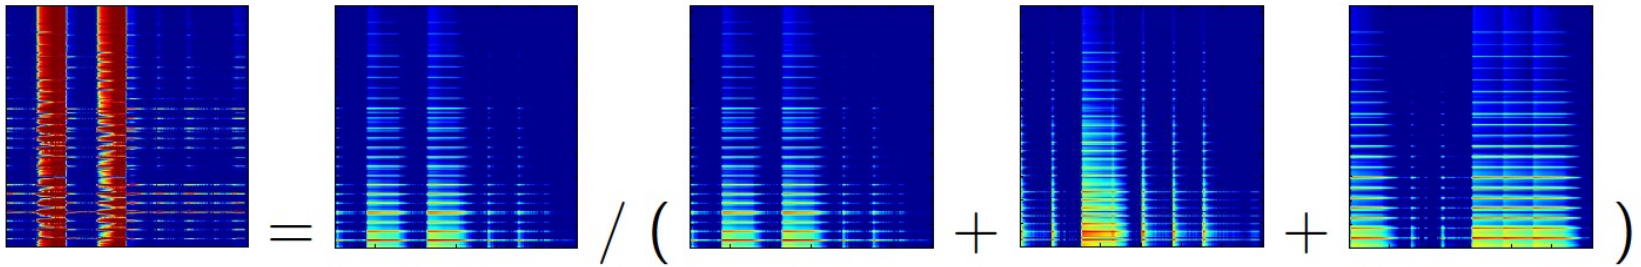

Source Separation: Traditional Approach

- To separate one component out of K:
  - 2. Masking on the input spectrogram to separate the component

$$|\hat{\mathbf{X}}_s| = \mathbf{M}_s \odot |\mathbf{X}|$$

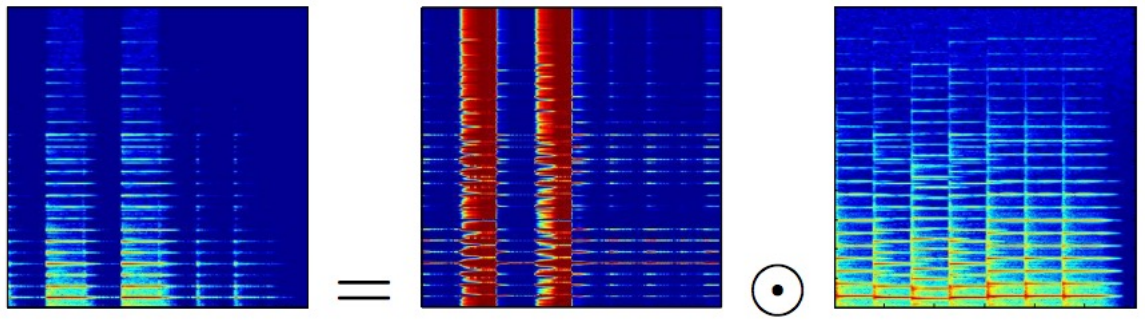

$|\hat{\mathbf{X}}_s|$ $\mathbf{M}_s$ $|\mathbf{X}|$

Our Ideas: Learning from Music Videos

❑ Internet music videos

- Keyword search without labeling
- >20 kinds of commonly seen musical instruments
- >1000 solos and duets

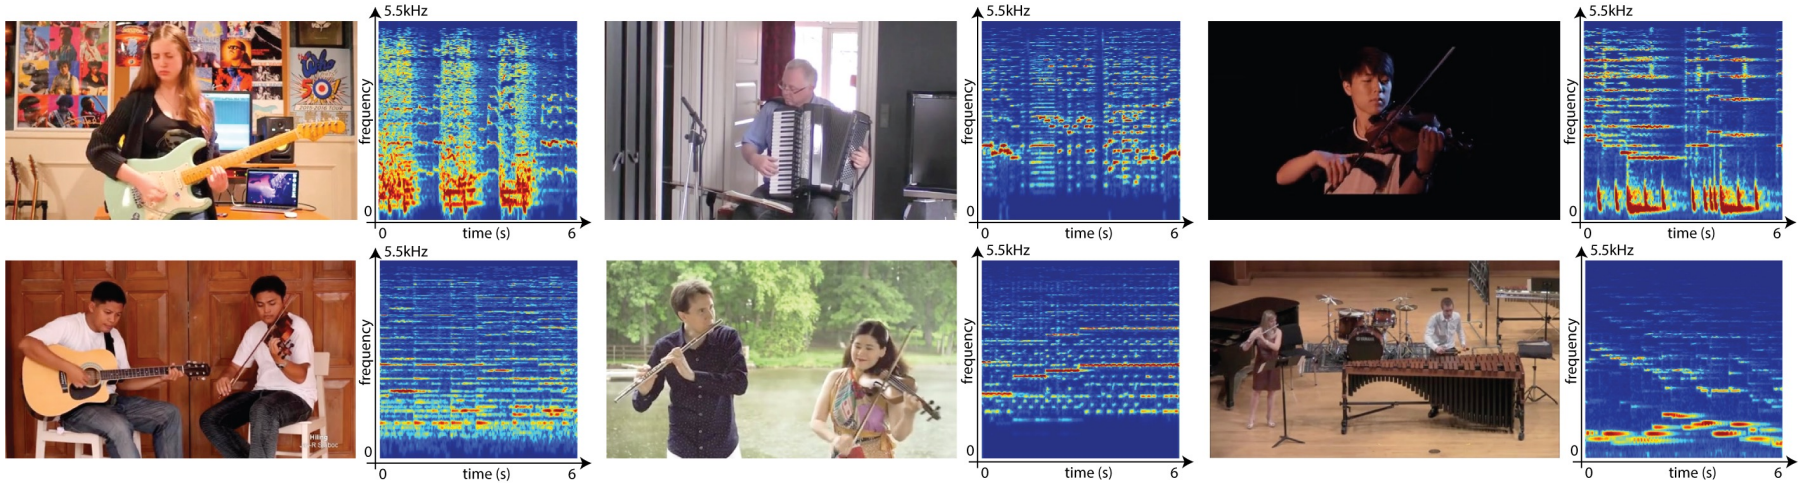

Zhao, Gan et al. "The Sound of Pixels." ECCV 2020.

Mixing the Sound

Zhao, Gan et al. “The Sound of Pixels.” ECCV 2020.

Vision to Rescue for Self-supervised Prediction

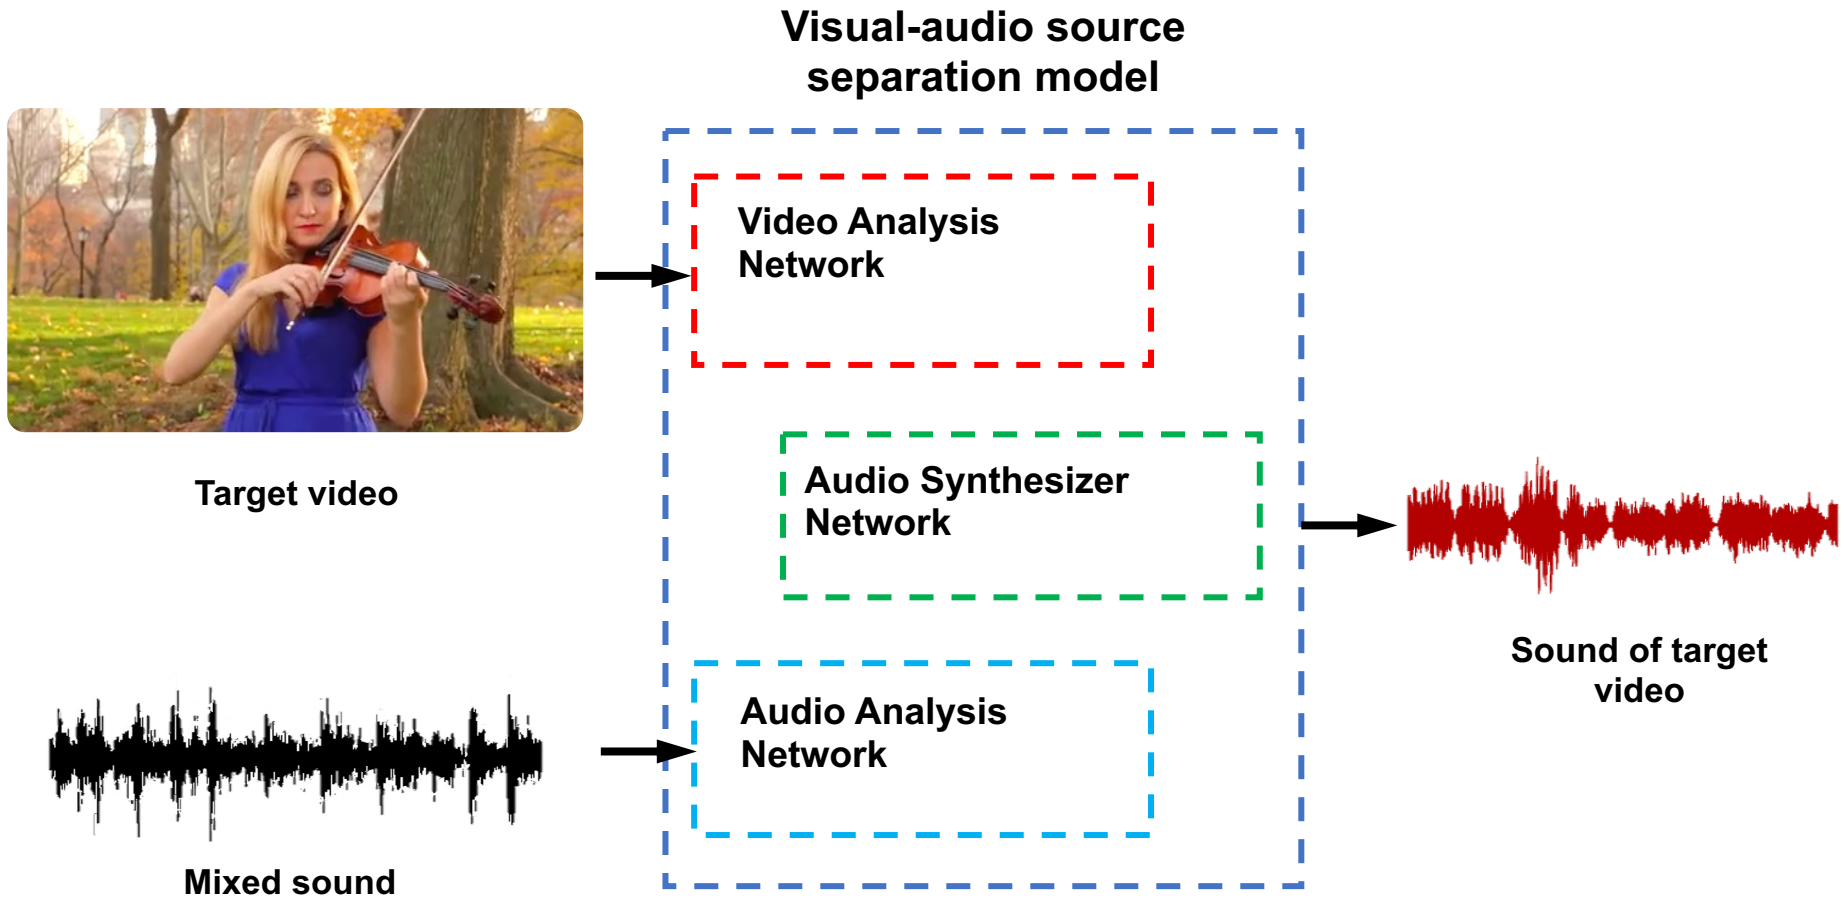

Mix-and-Separate Framework

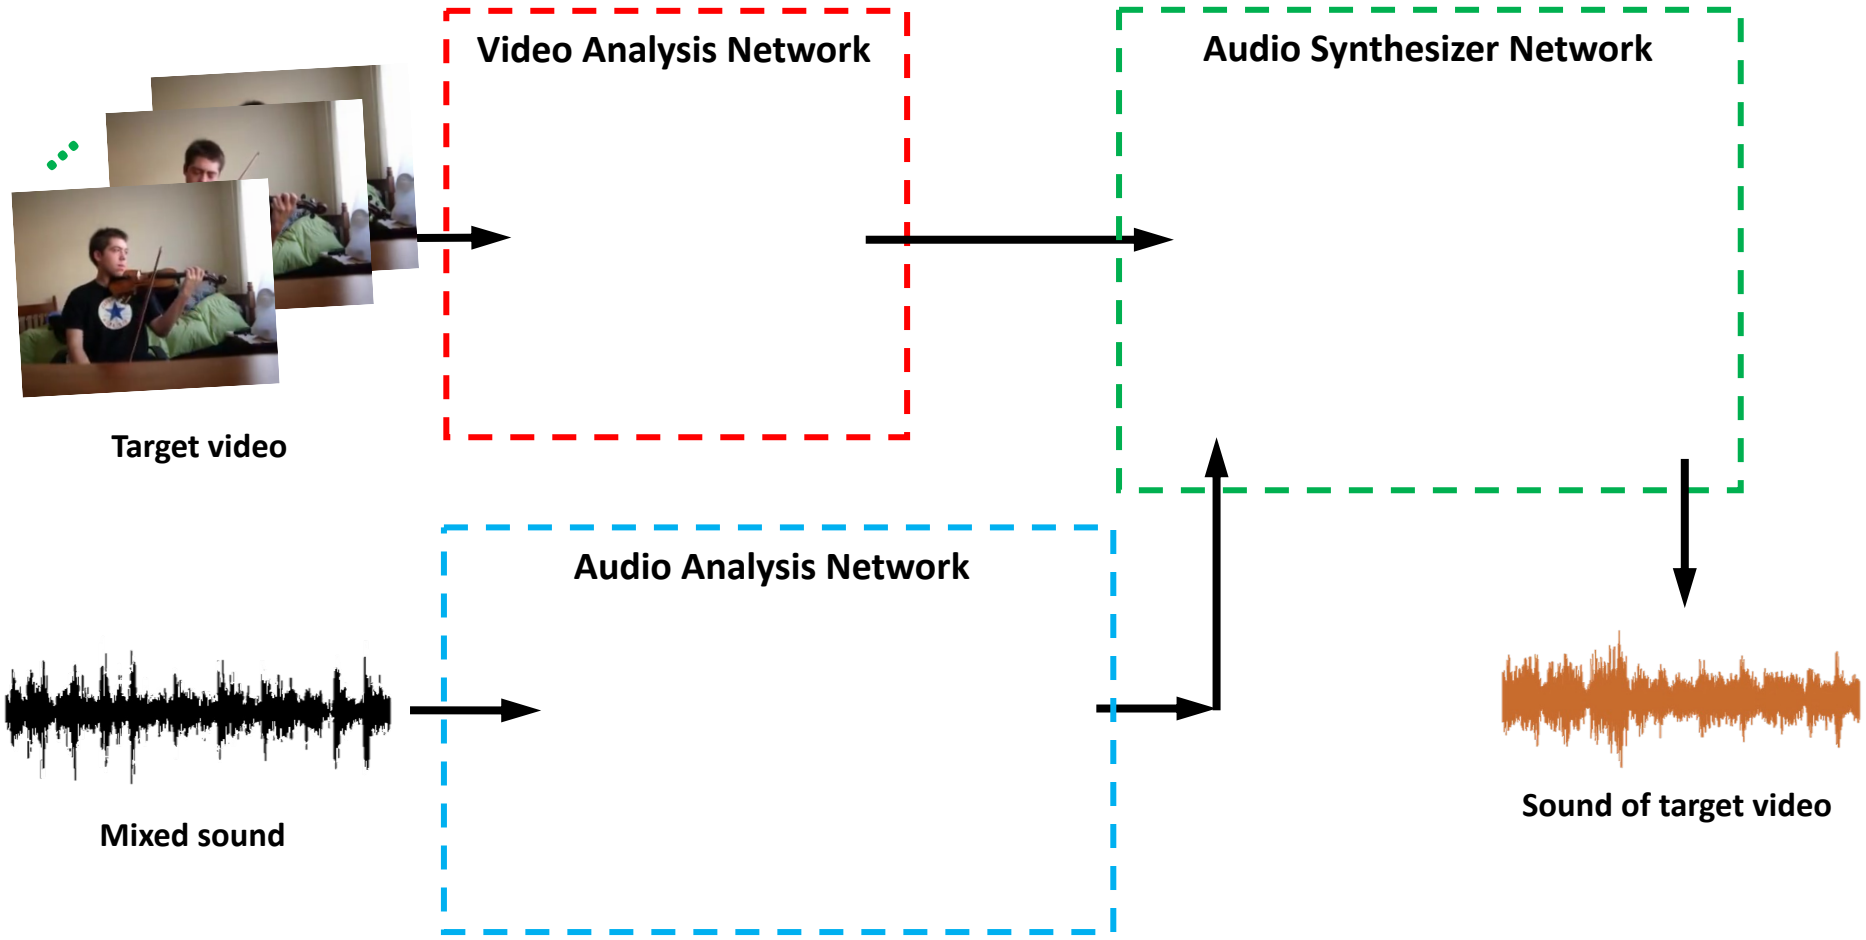

Mix-and-Separate Framework

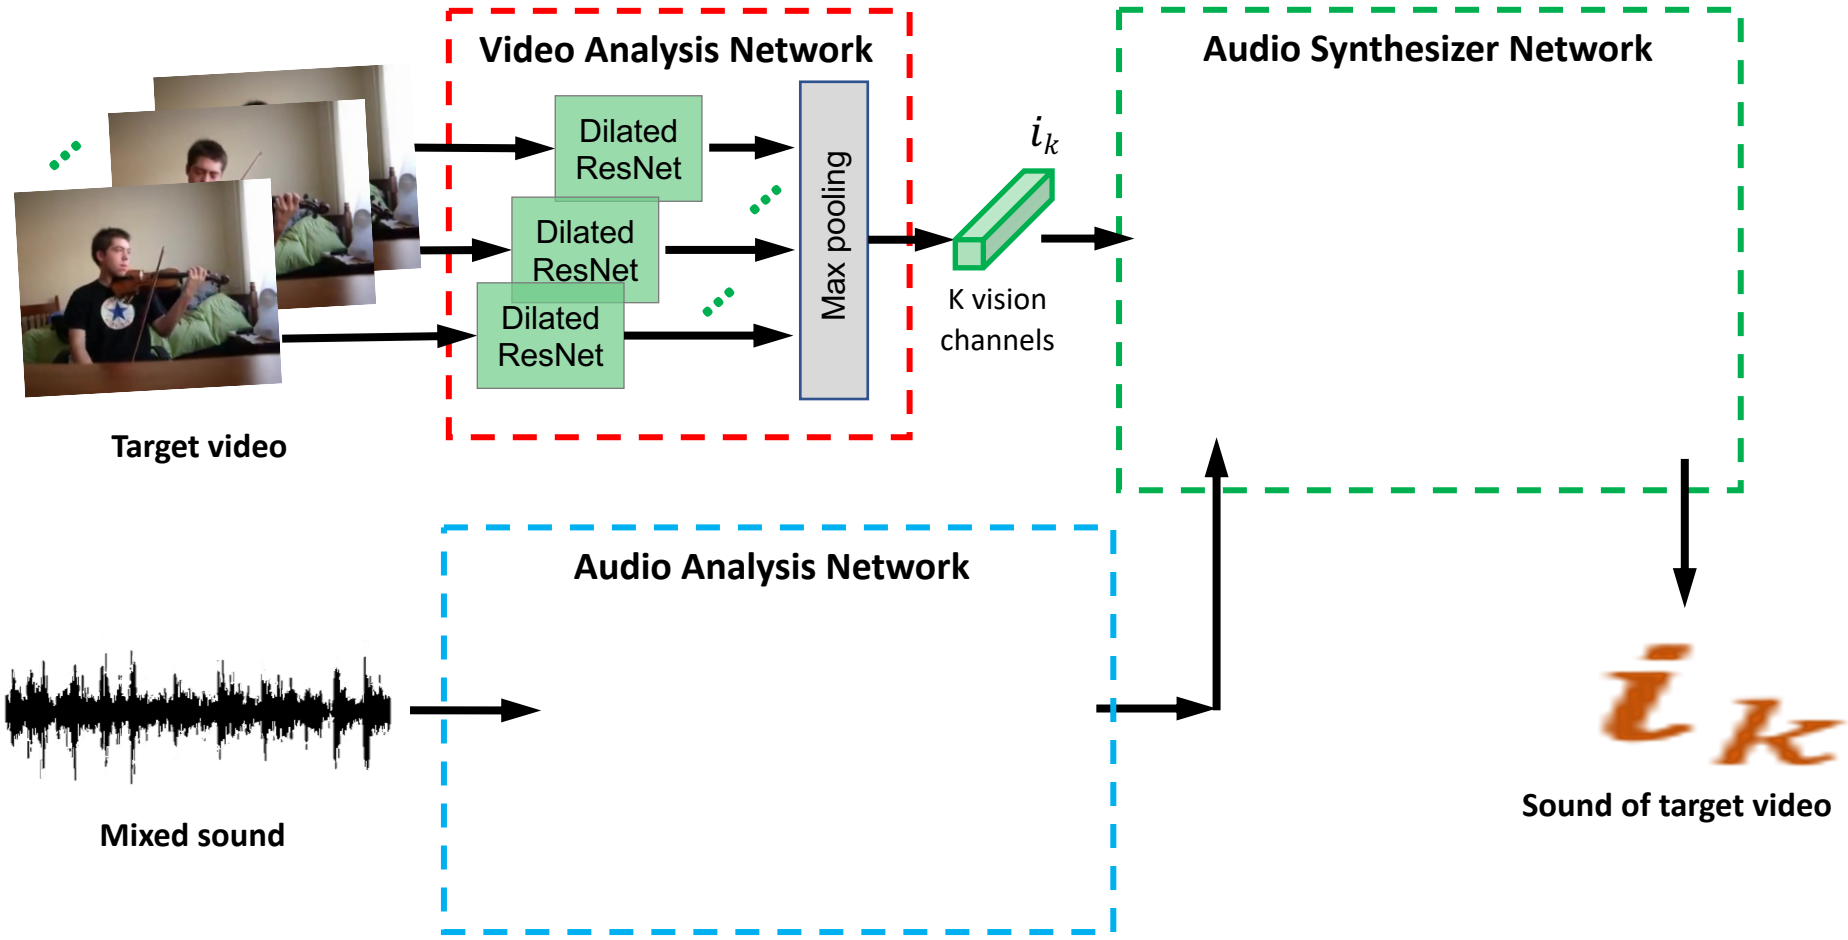

Mix-and-Separate Framework

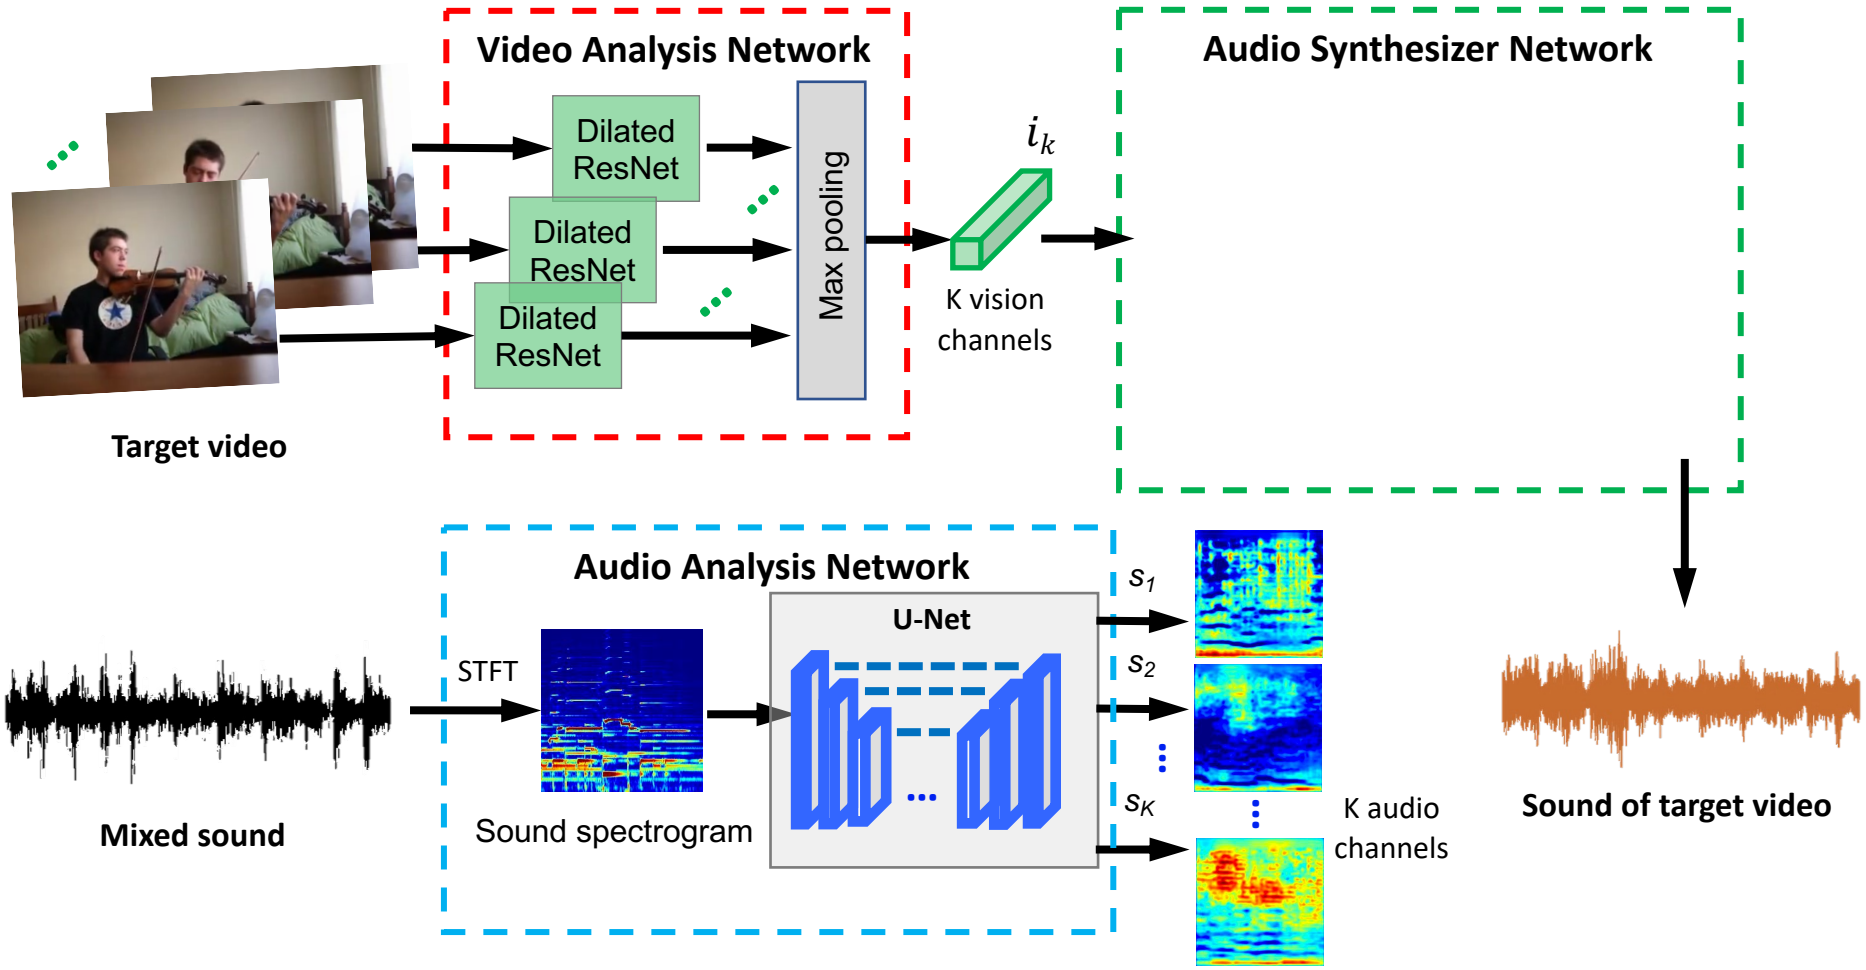

Mix-and-Separate Framework

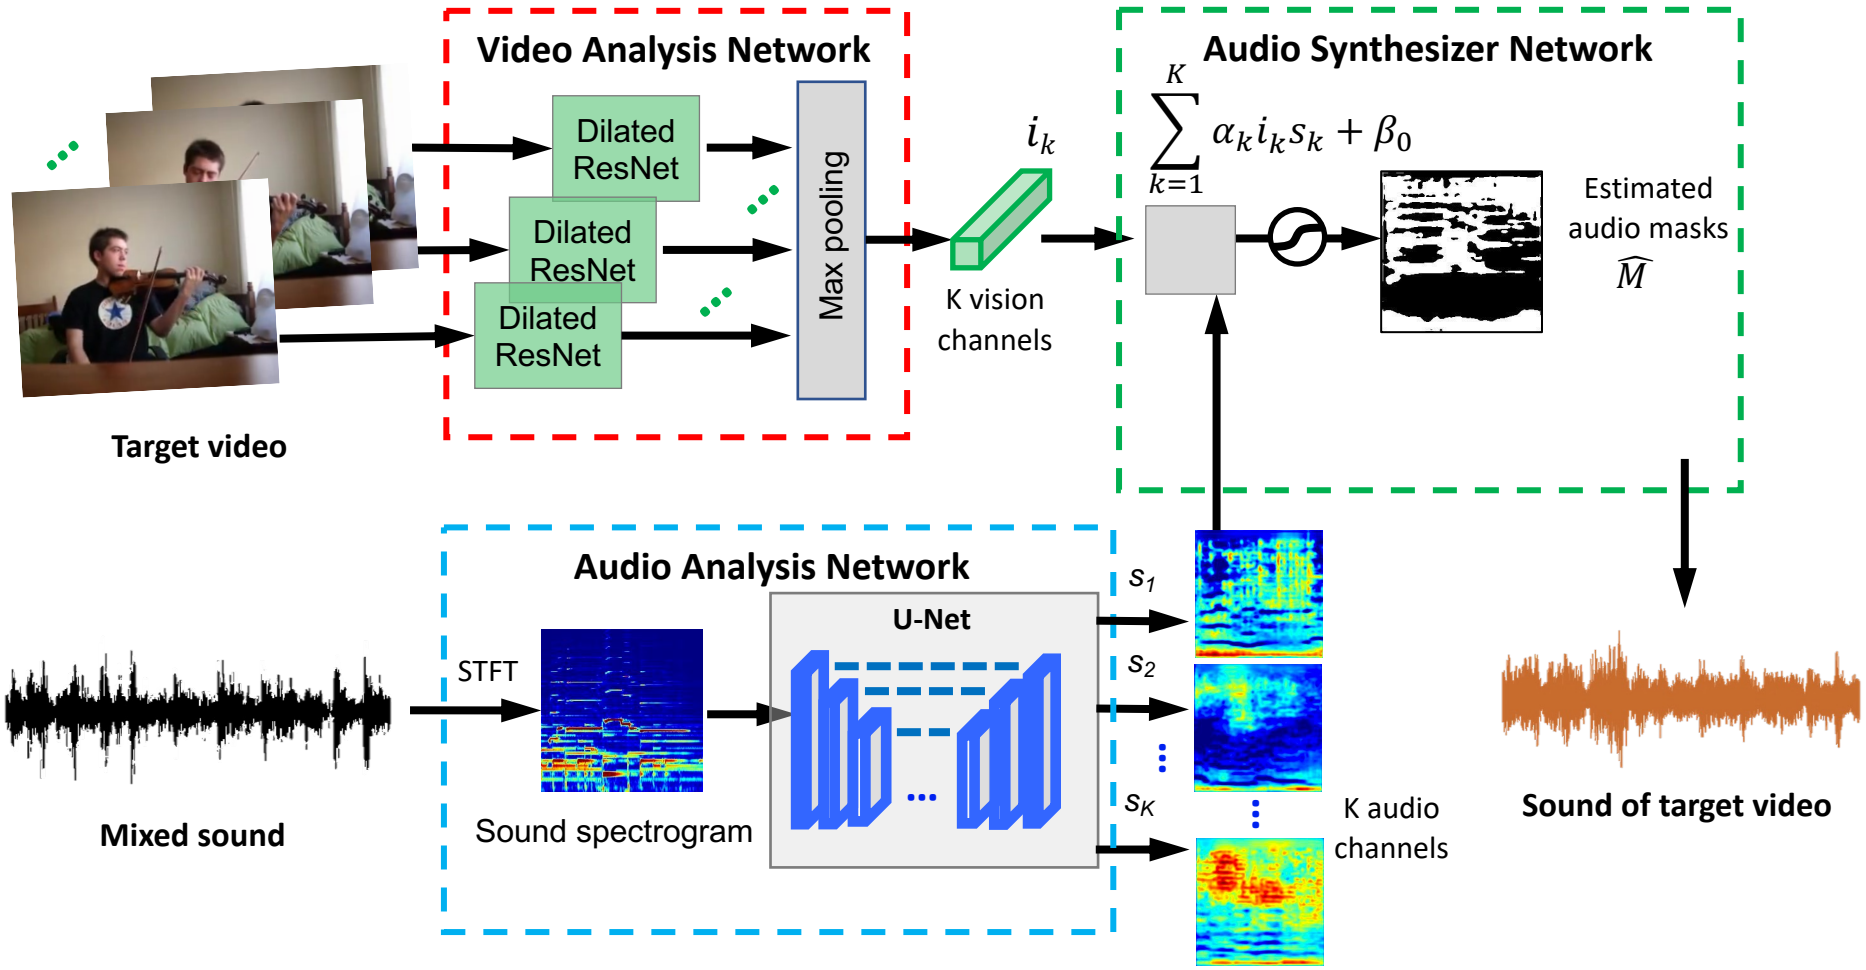

Mix-and-Separate Framework

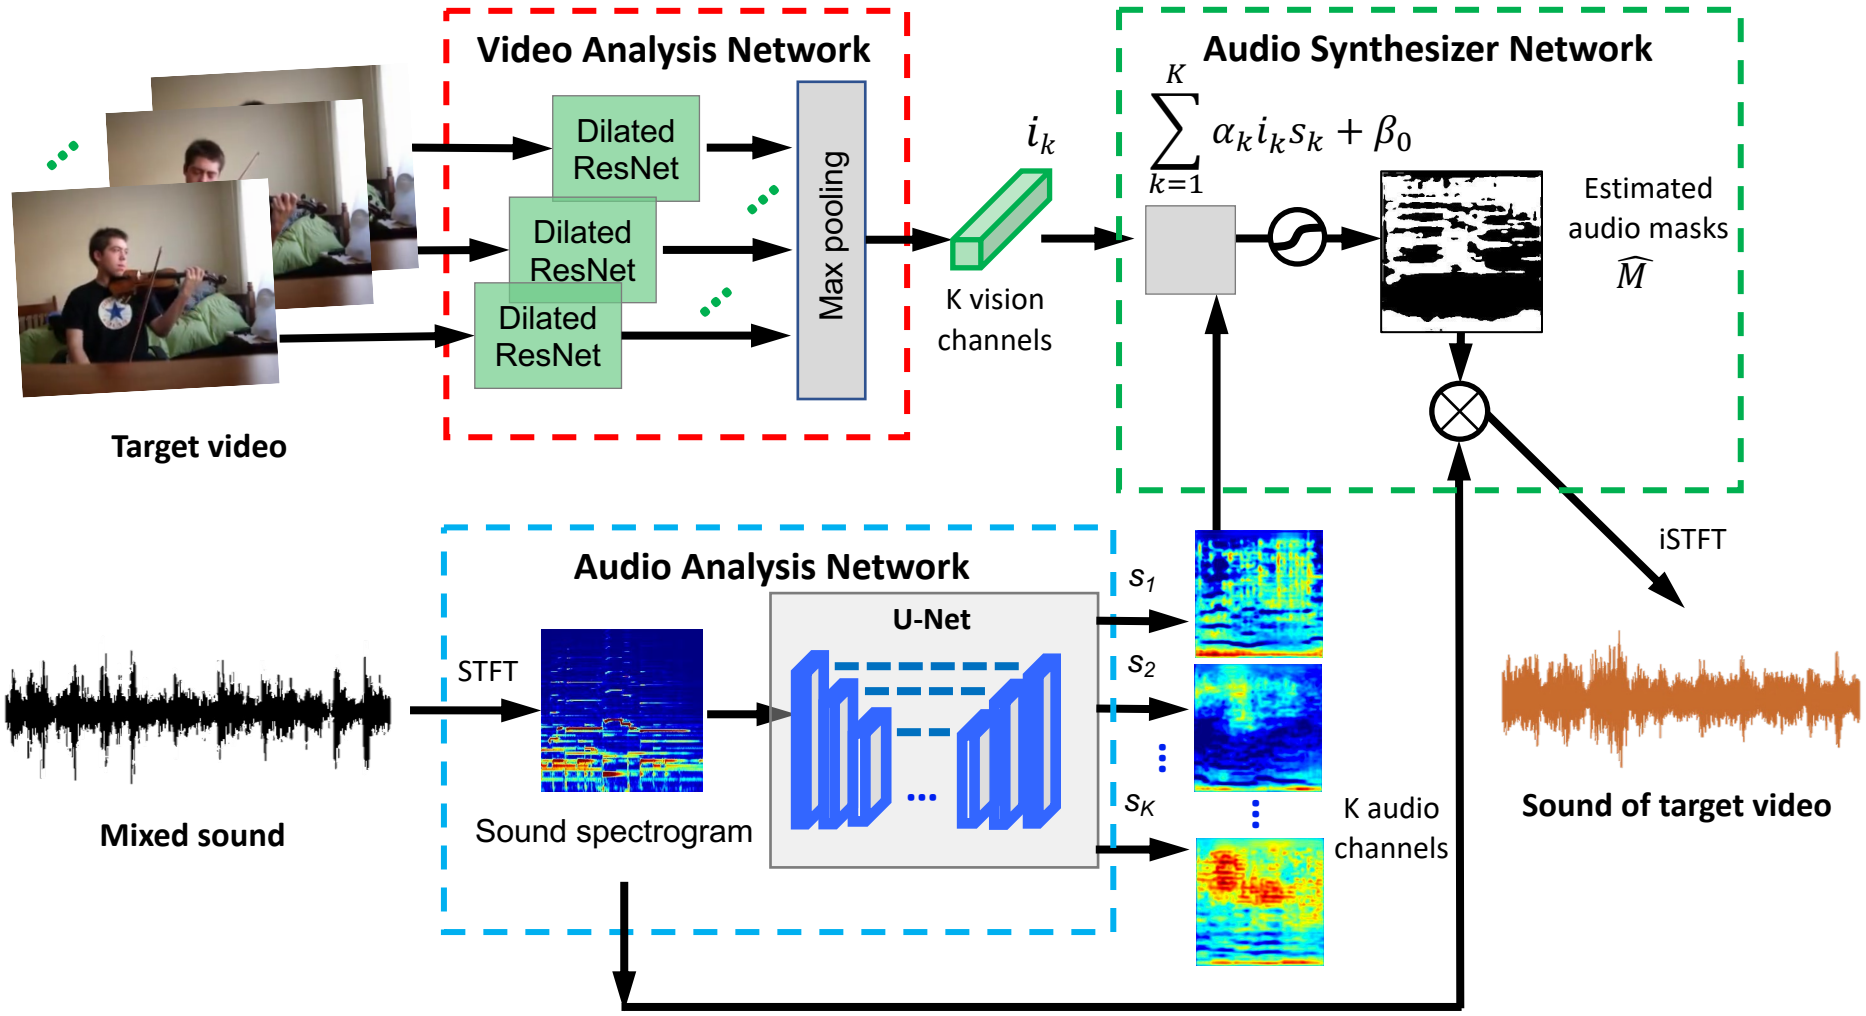

Test Time

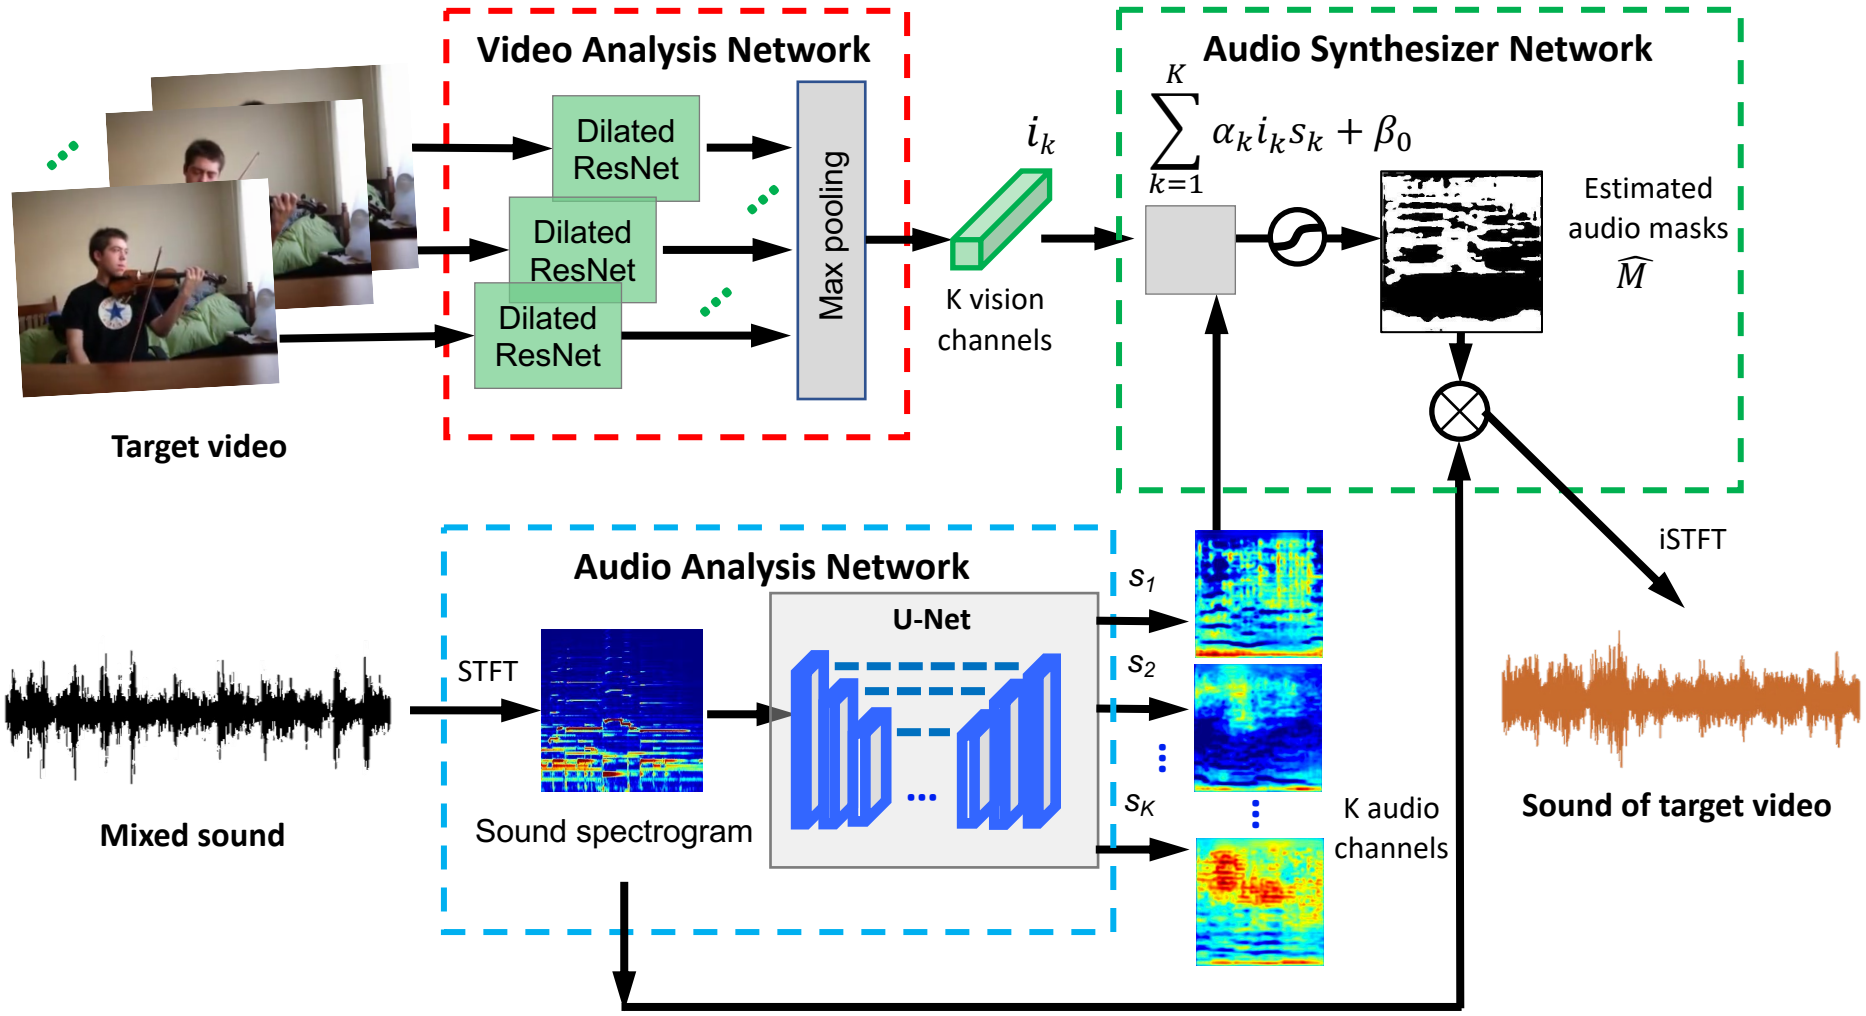
Video Analysis Network
Dilated ResNet
Dilated ResNet
Dilated ResNet
Max pooling
i_k
K vision channels
Target video
Audio Synthesizer Network
$\sum_{k=1}^{K} \alpha_k i_k s_k + \beta_0$
Estimated audio masks
$\widehat{M}$
iSTFT
Audio Analysis Network
U-Net
STFT
Sound spectrogram
Mixed sound
s_1
s_2
s_K
K audio channels
Sound of target video

Test Time: using Pixel Feature instead

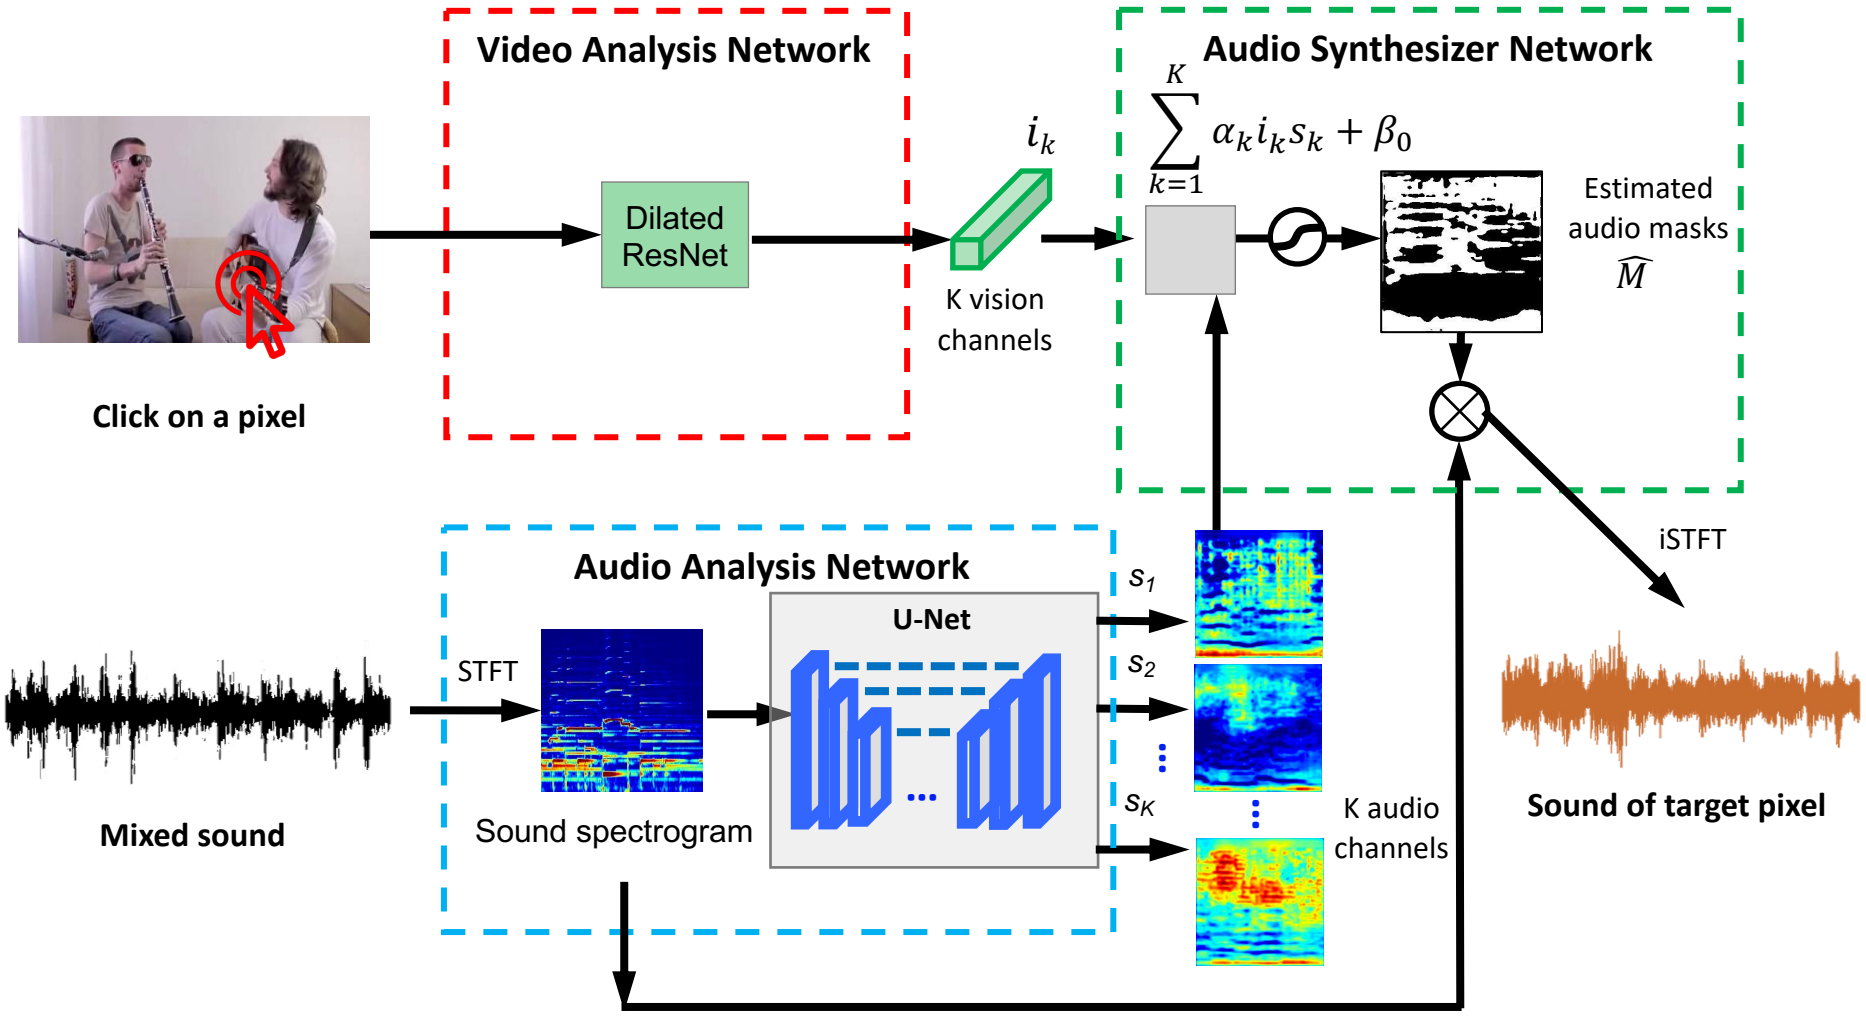

Original Video

The sound of clicked object…

The sound of clicked object…

The sound of clicked object…

Original Video

The sound of clicked object…

The sound of clicked object…

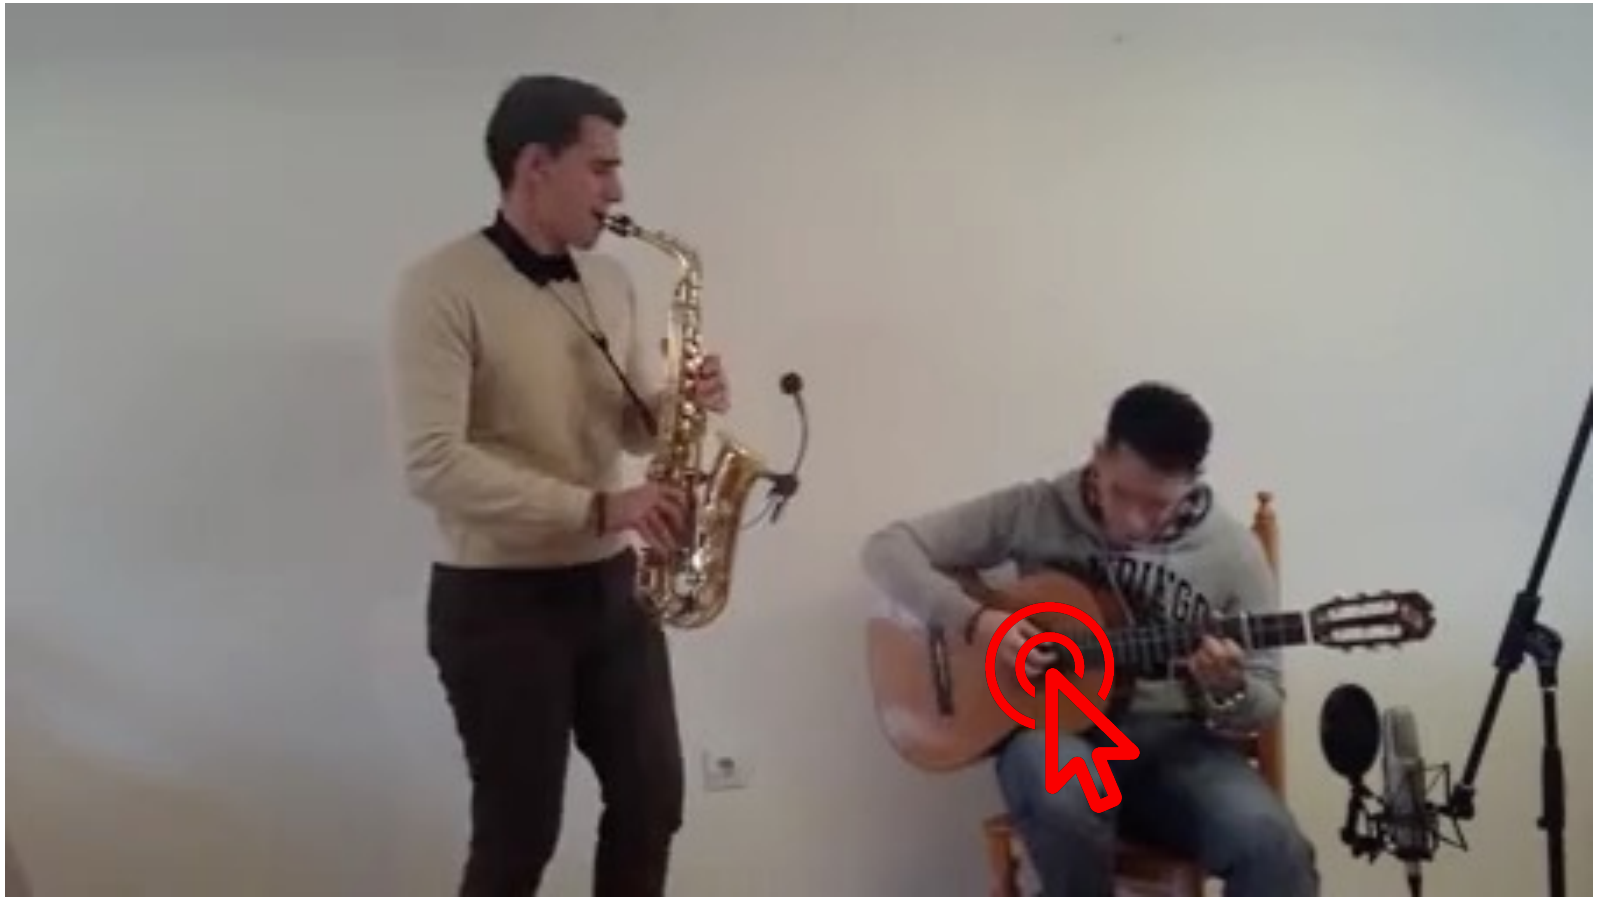

The sound of clicked object…

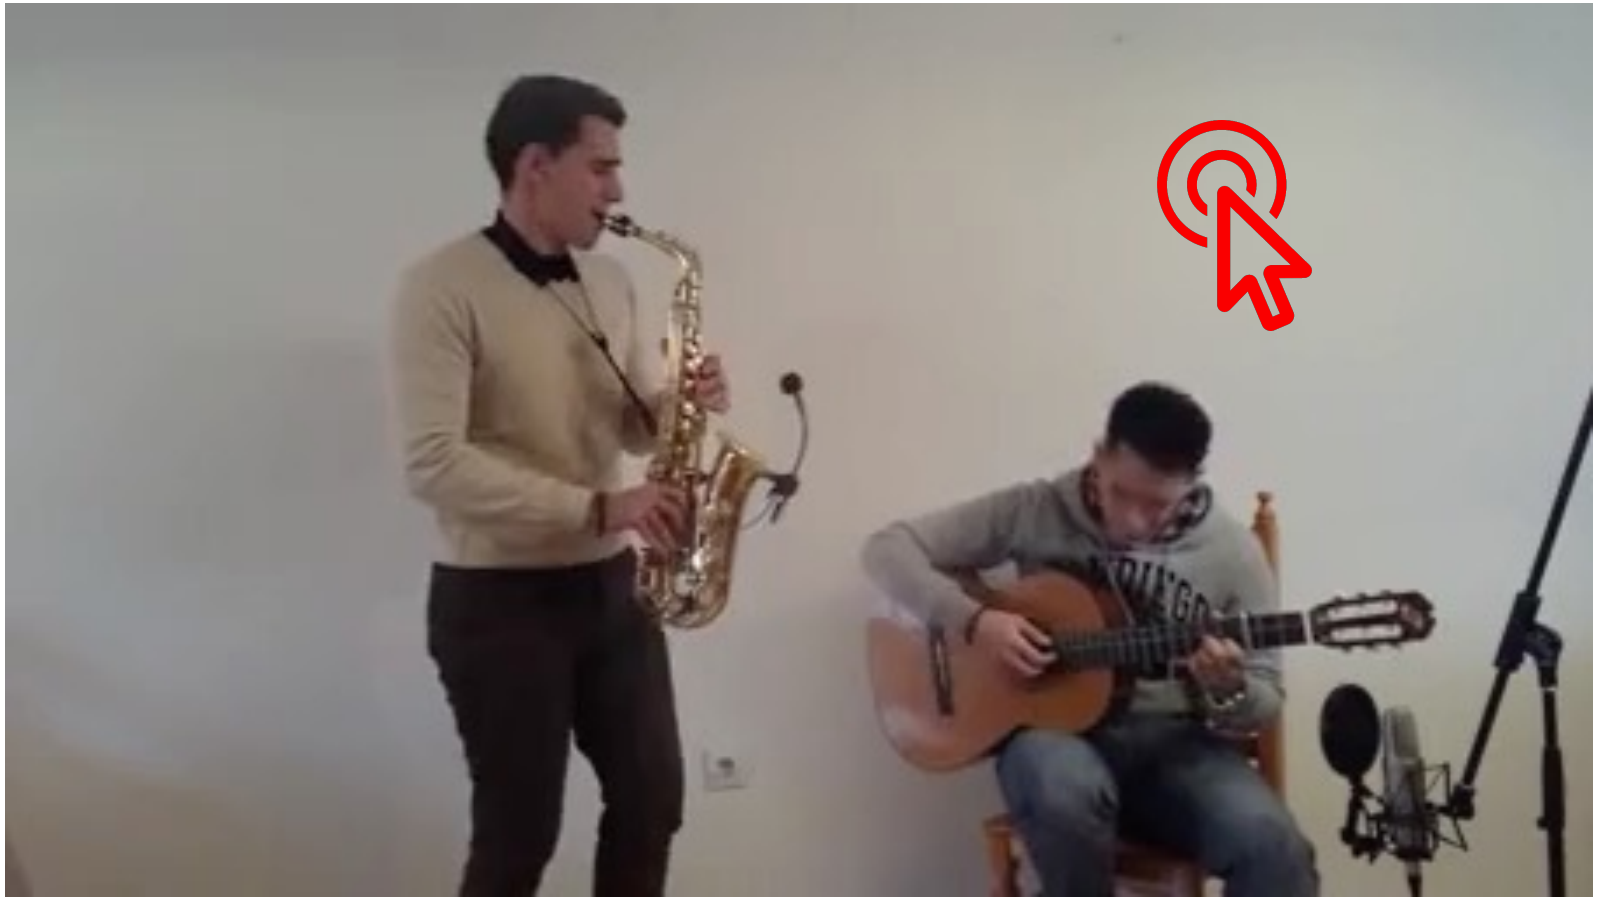

Application: Music Remix

Application: Music Remix

Limitation

- Most existing methods use **raw pixel** or **optical flow** as input.

- Problem: limited to **separate multiple instruments of the same types**.

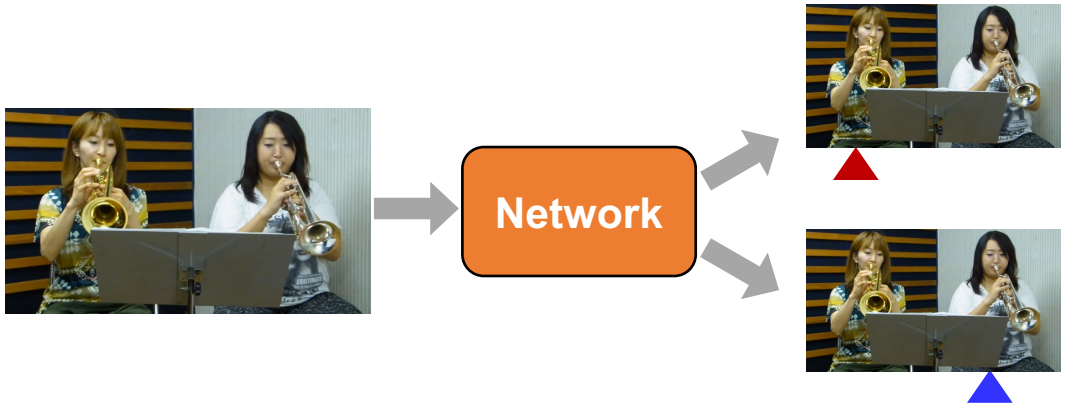

Our Ideas

❑ Problem: limited to **separate multiple instruments of the same types**.

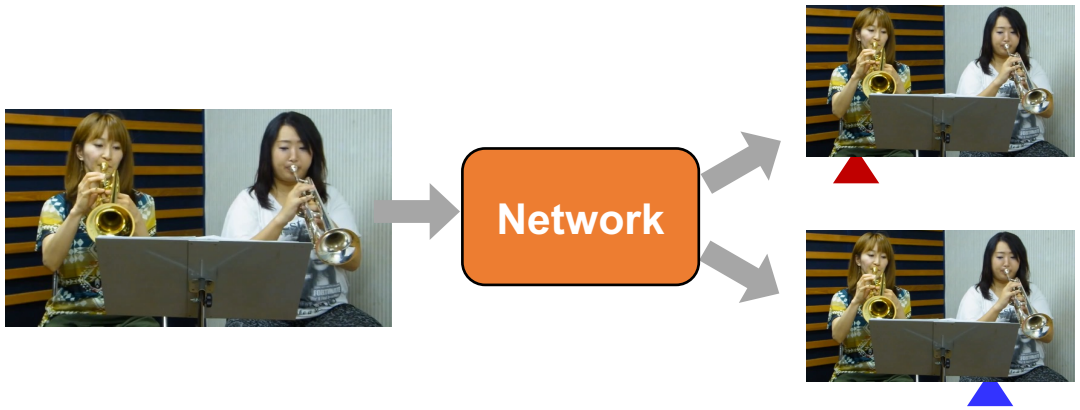

❑ We propose to Identifying a melody by studying a musician's body language using ``**Music Gesture**".

Gan et al. "Music Gesture for Visual Sound Separation." CVPR 2020.

Music Gesture

- Keypoint-based structured representations

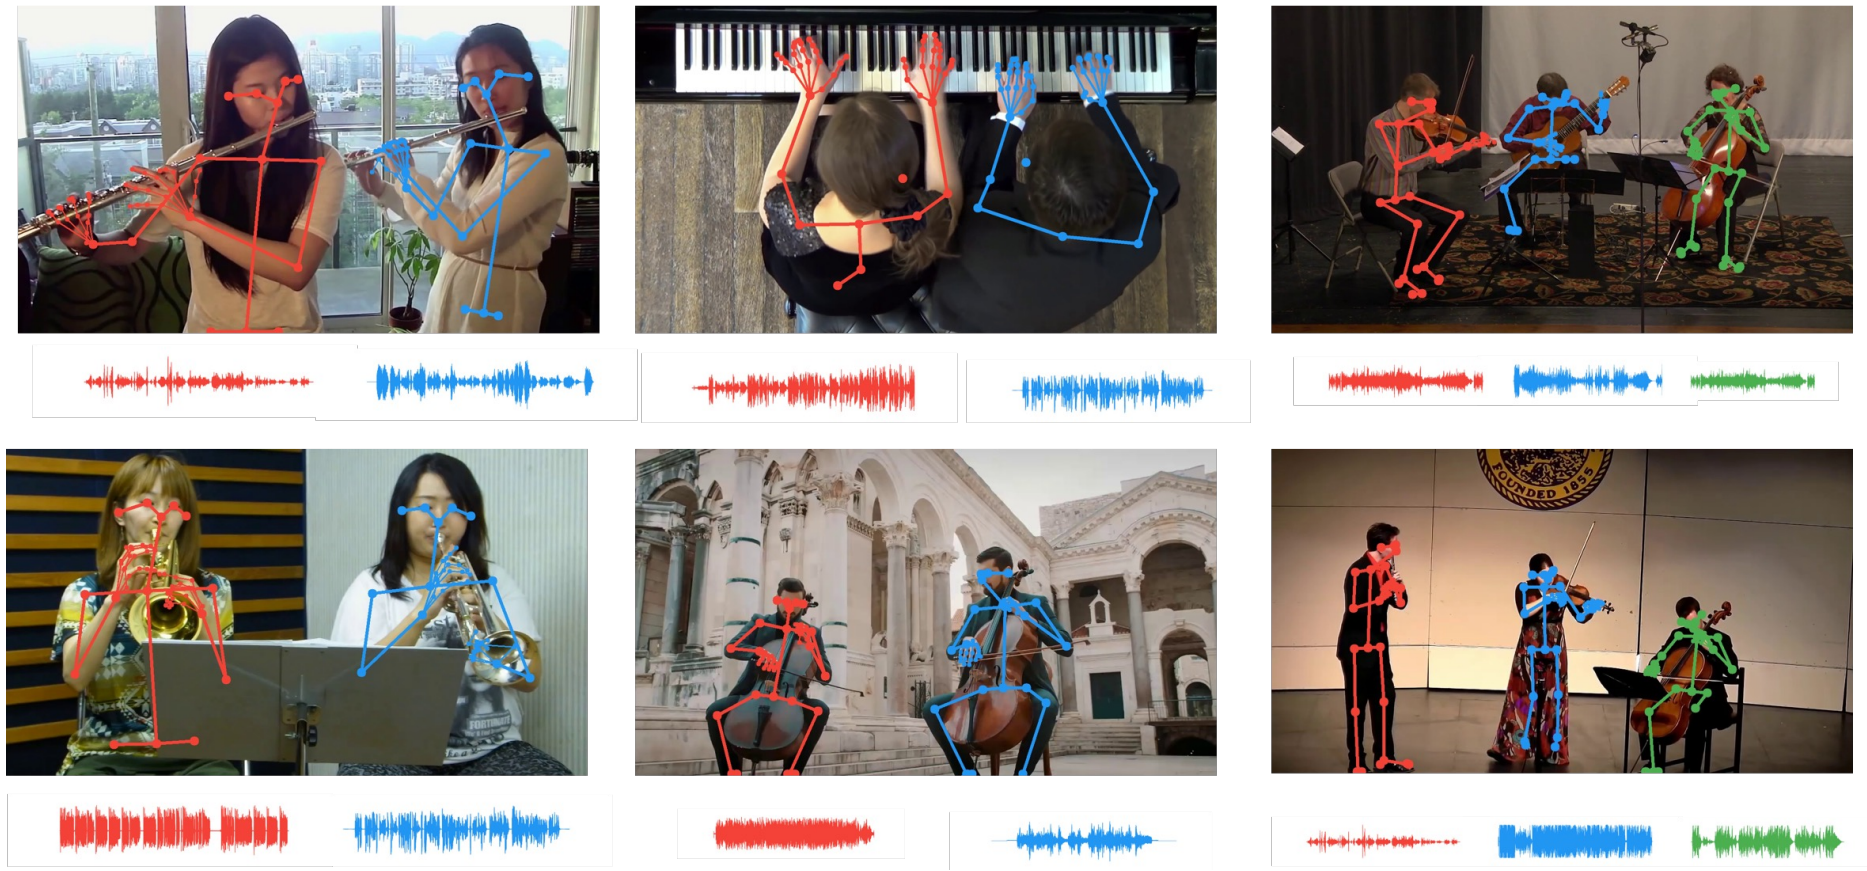

Visual sound separation results

Sound of Motion

Mixed sound

Separated sound1

Separated sound2

Music Gesture

Mixed sound

Separated sound1

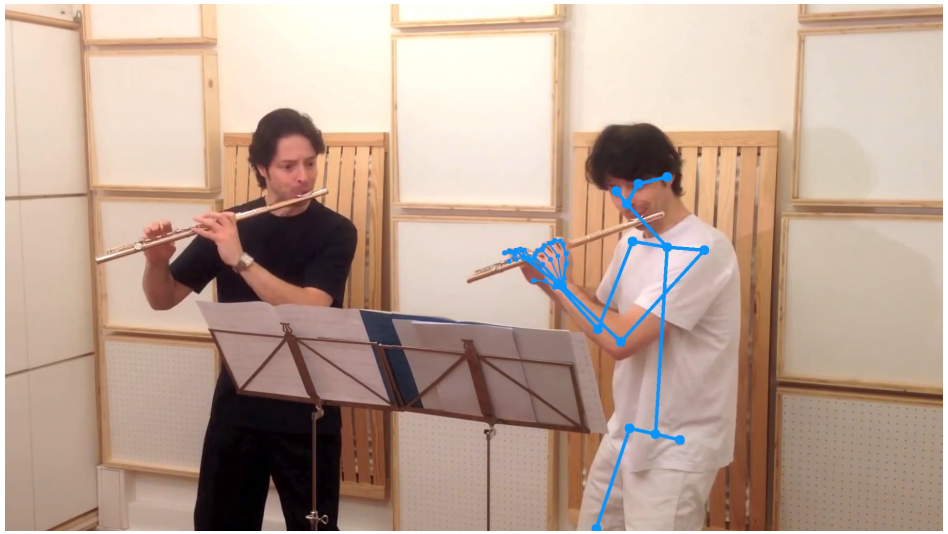

Separated sound2

Sound of Motion

Mixed sound

Separated sound1

Separated sound2

Music Gesture

Mixed sound

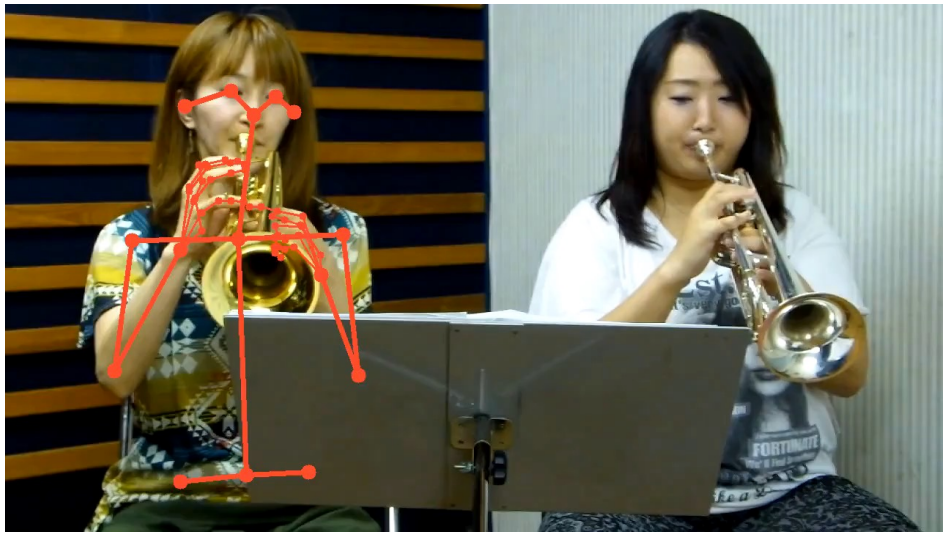

Separated sound1

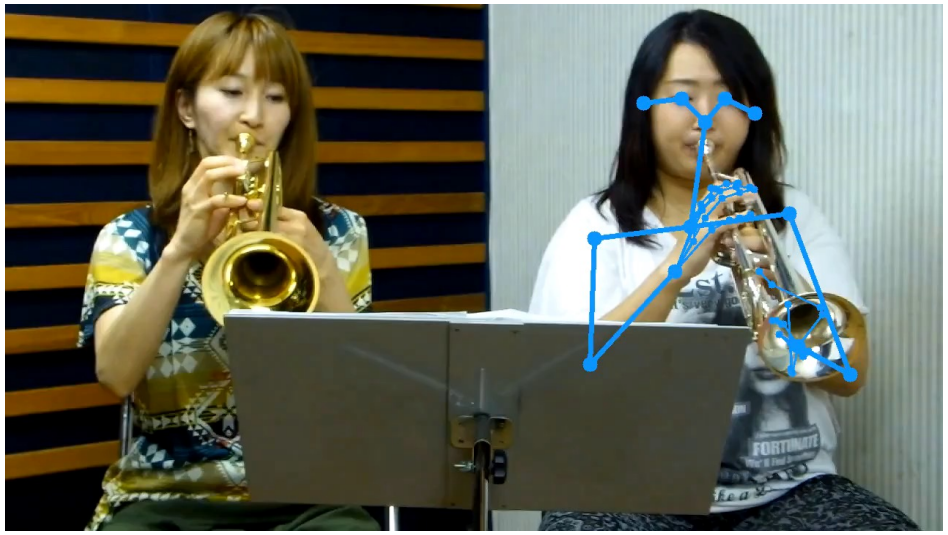

Separated sound2

Sound of Motion

Mixed sound

Separated sound1

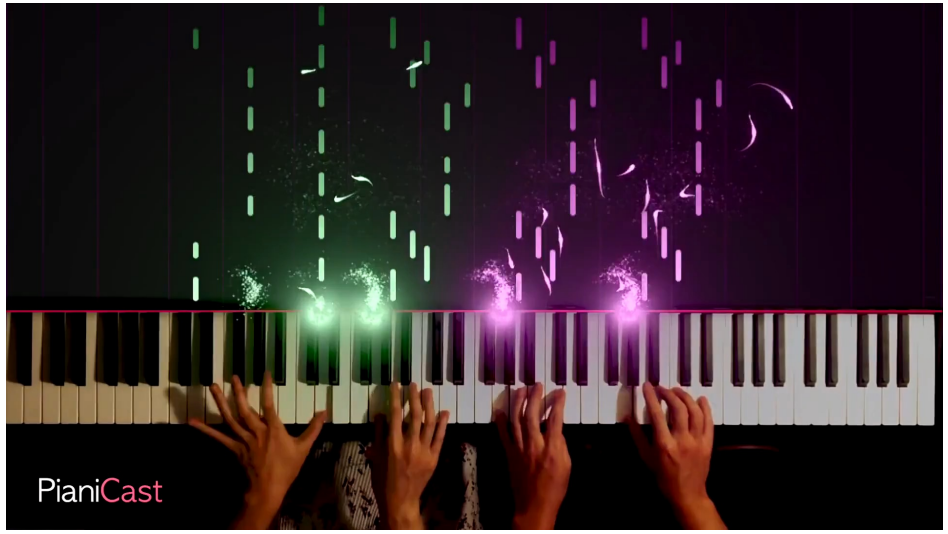

Separated sound2

Music Gesture

Mixed sound

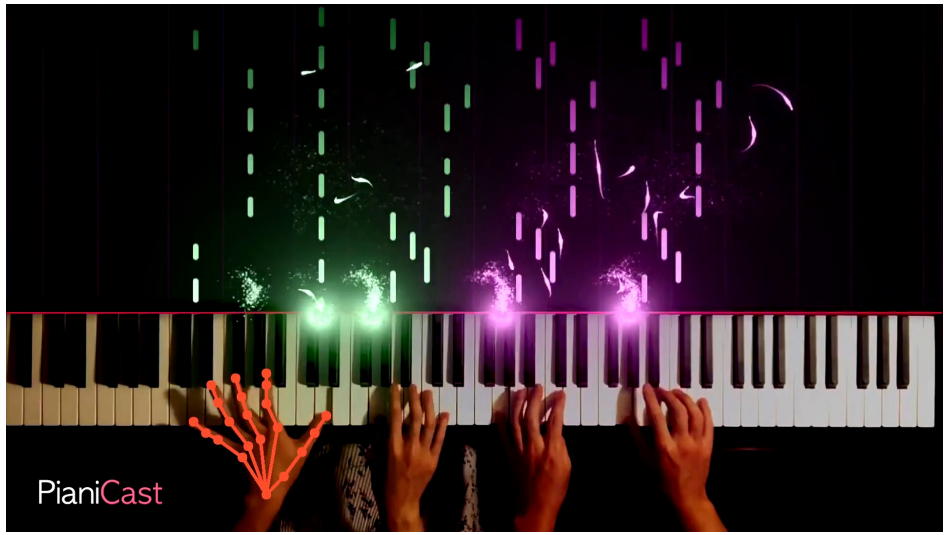

Separated sound1

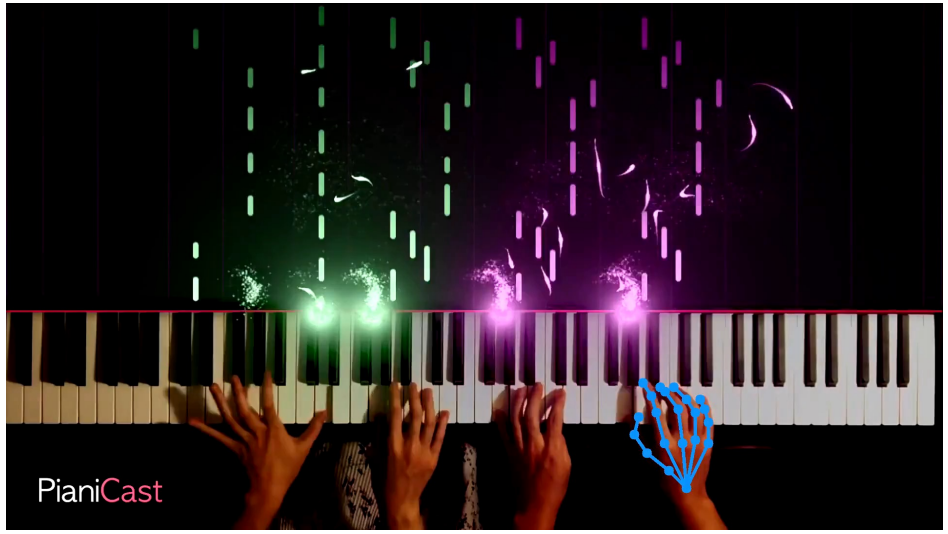

Separated sound2

Multiple instruments

Music Gesture

Mixed sound

Separated sound1

Separated sound2

Music Gesture

Separated sound3

Mixed sound

Separated sound4

Attention Map of Key points

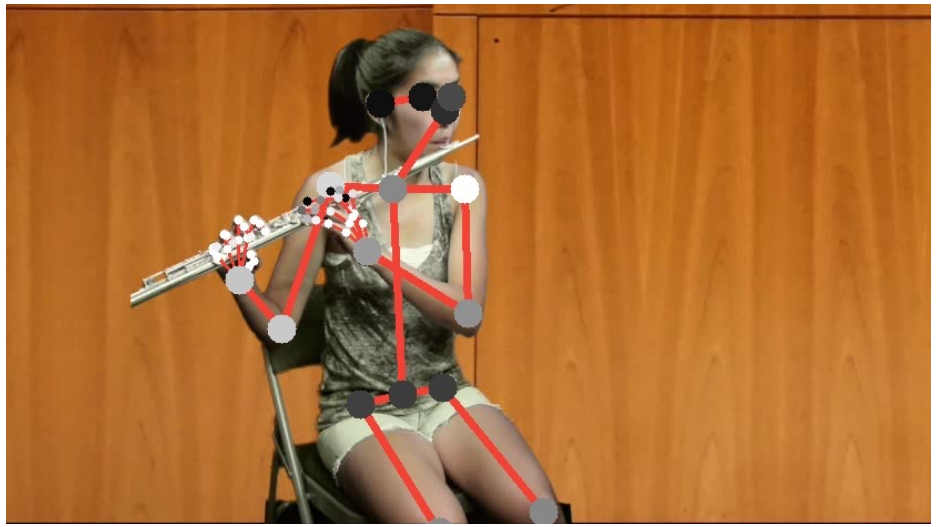

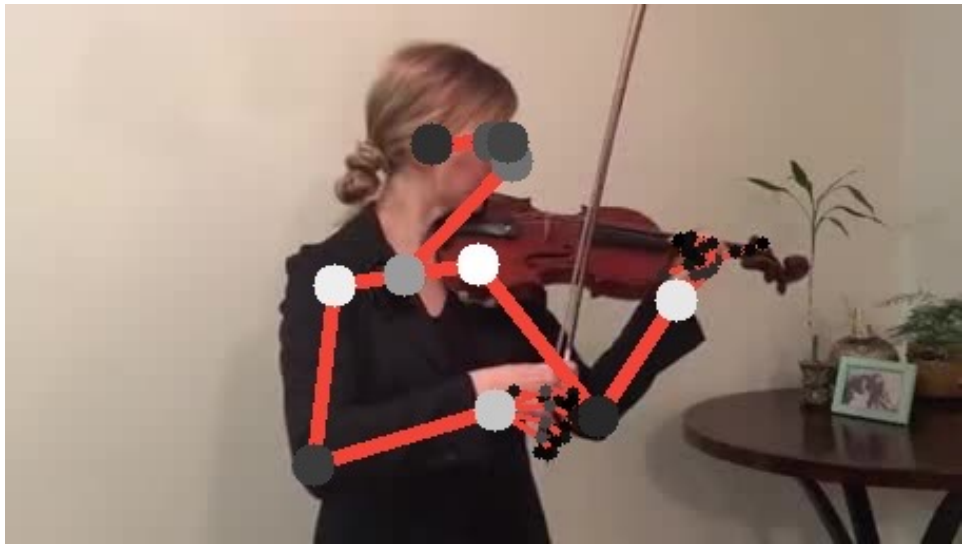

The sound of body parts

Mixed sound

Separated sound1

Separated sound2

Can we generate music from videos?

Given a silent music performance video…

Silent music performance video

Gan et al. “Foley Music: Learning to Generate Music from Videos.” ECCV 2020.

Can we generate music from videos?

…we aim to generate plausible music.

Deep Neural Network

Silent music performance video

Performance with generated sound

Gan et al. "Foley Music: Learning to Generate Music from Videos." ECCV 2020.

Challenges

- Hard to learn visual-audio mappings from unlabeled video
- Three things matter：
  - Visual perception module → interactions between instrument and player
  - Audio representation → musical rules, easy to predict from visual signals
  - Visual-audio model → association between two modalities

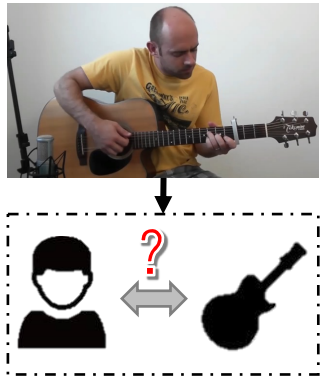

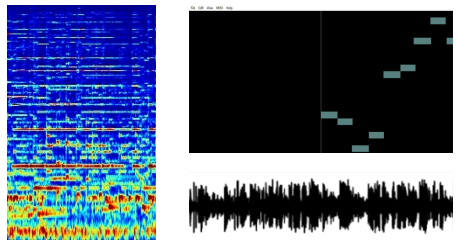

Choose ?

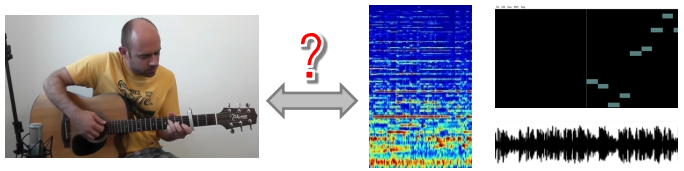

Challenges

- Hard to learn visual-audio mappings from unlabeled video
- Three things matter：
  - Visual perception module → interactions between instrument and player
  - Audio representation → musical rules, easy to predict from visual signals
  - Visual-audio model → association between two modalities

We use **body keypoints** to explicitly model the body and finger.

We use **Musical Instrument Digital Interface (MIDI)** to represent music.

MIDI Event Representations

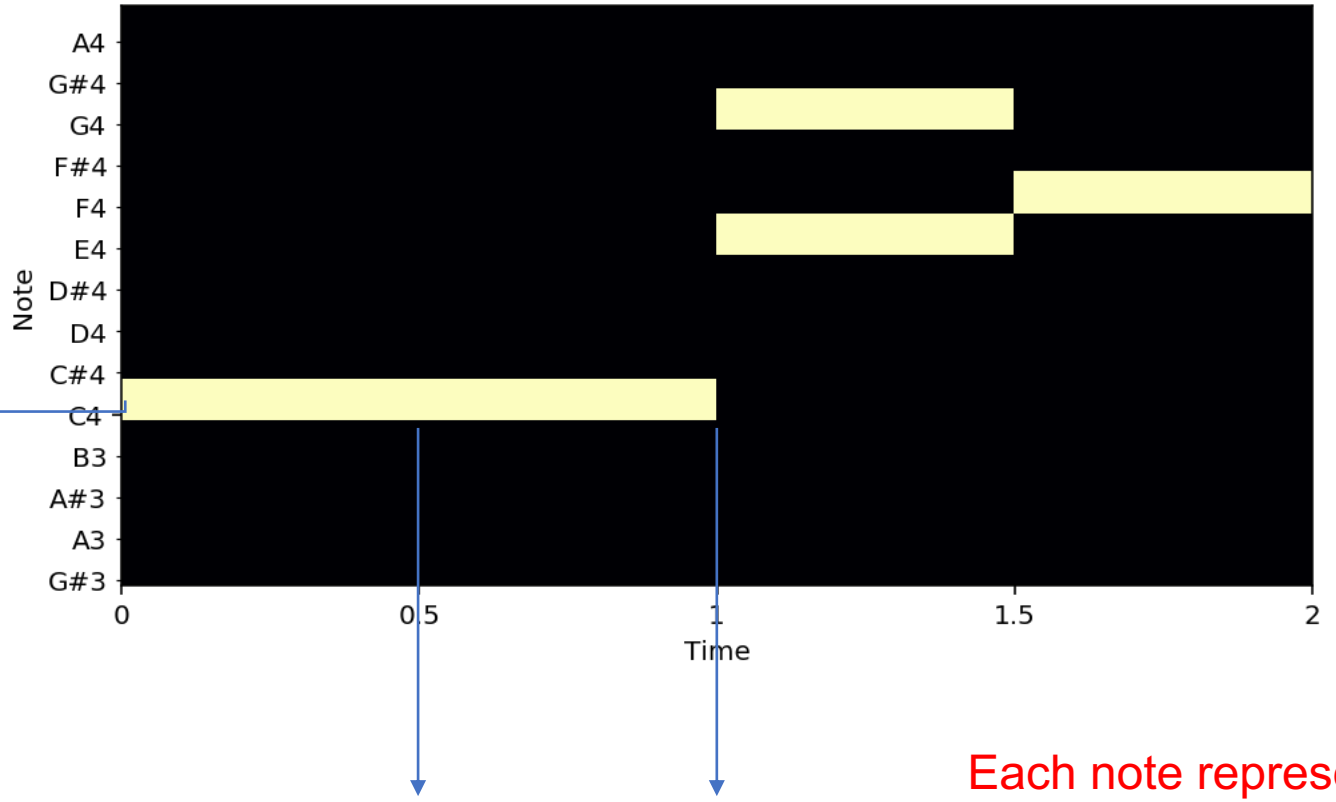

Velocity Event, Note On Event	Time Shift Event	Note Off Event
201	39	237	127

Each note represented as a sequence of MIDI Events

- Note On Event$\in [0, 88]$, based on Pitch
- Note Off Event$\in [88, 176]$, based on Pitch
- Velocity Event$\in [176, 208]$, based on Velocity
- Time Shift Event$\in [208, 240]$, based on Duration

Our method

We first encode human keypoints by using graph CNN.

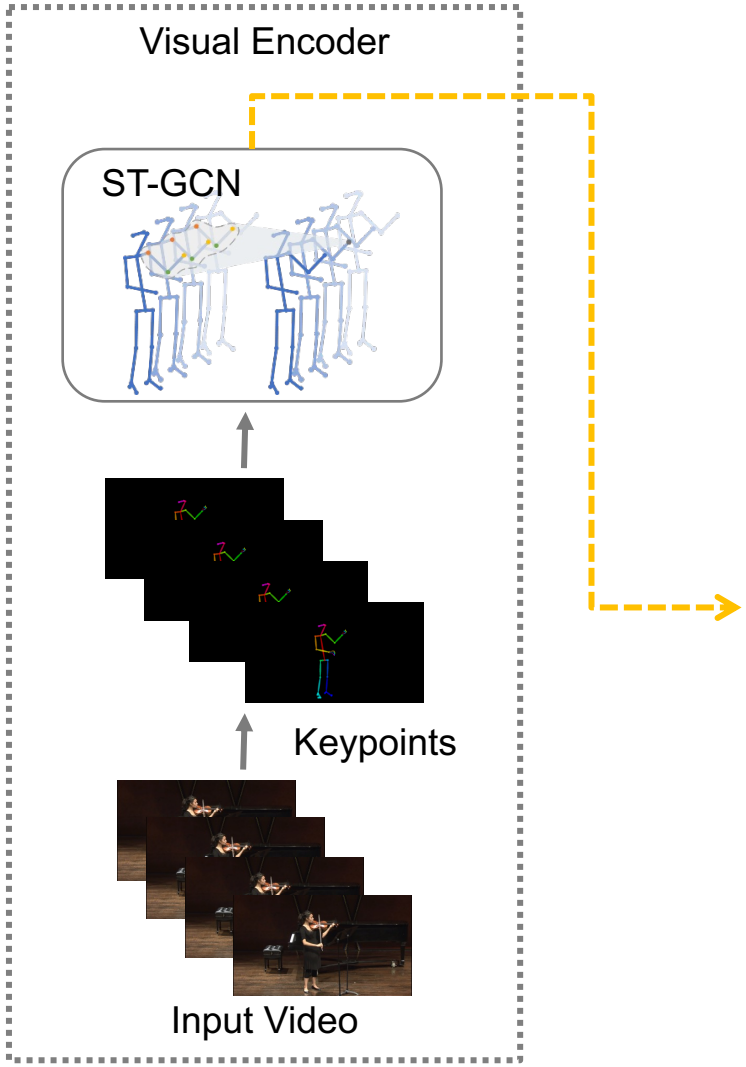

Our method

We then translate the keypoints features into MIDI using Transformers.

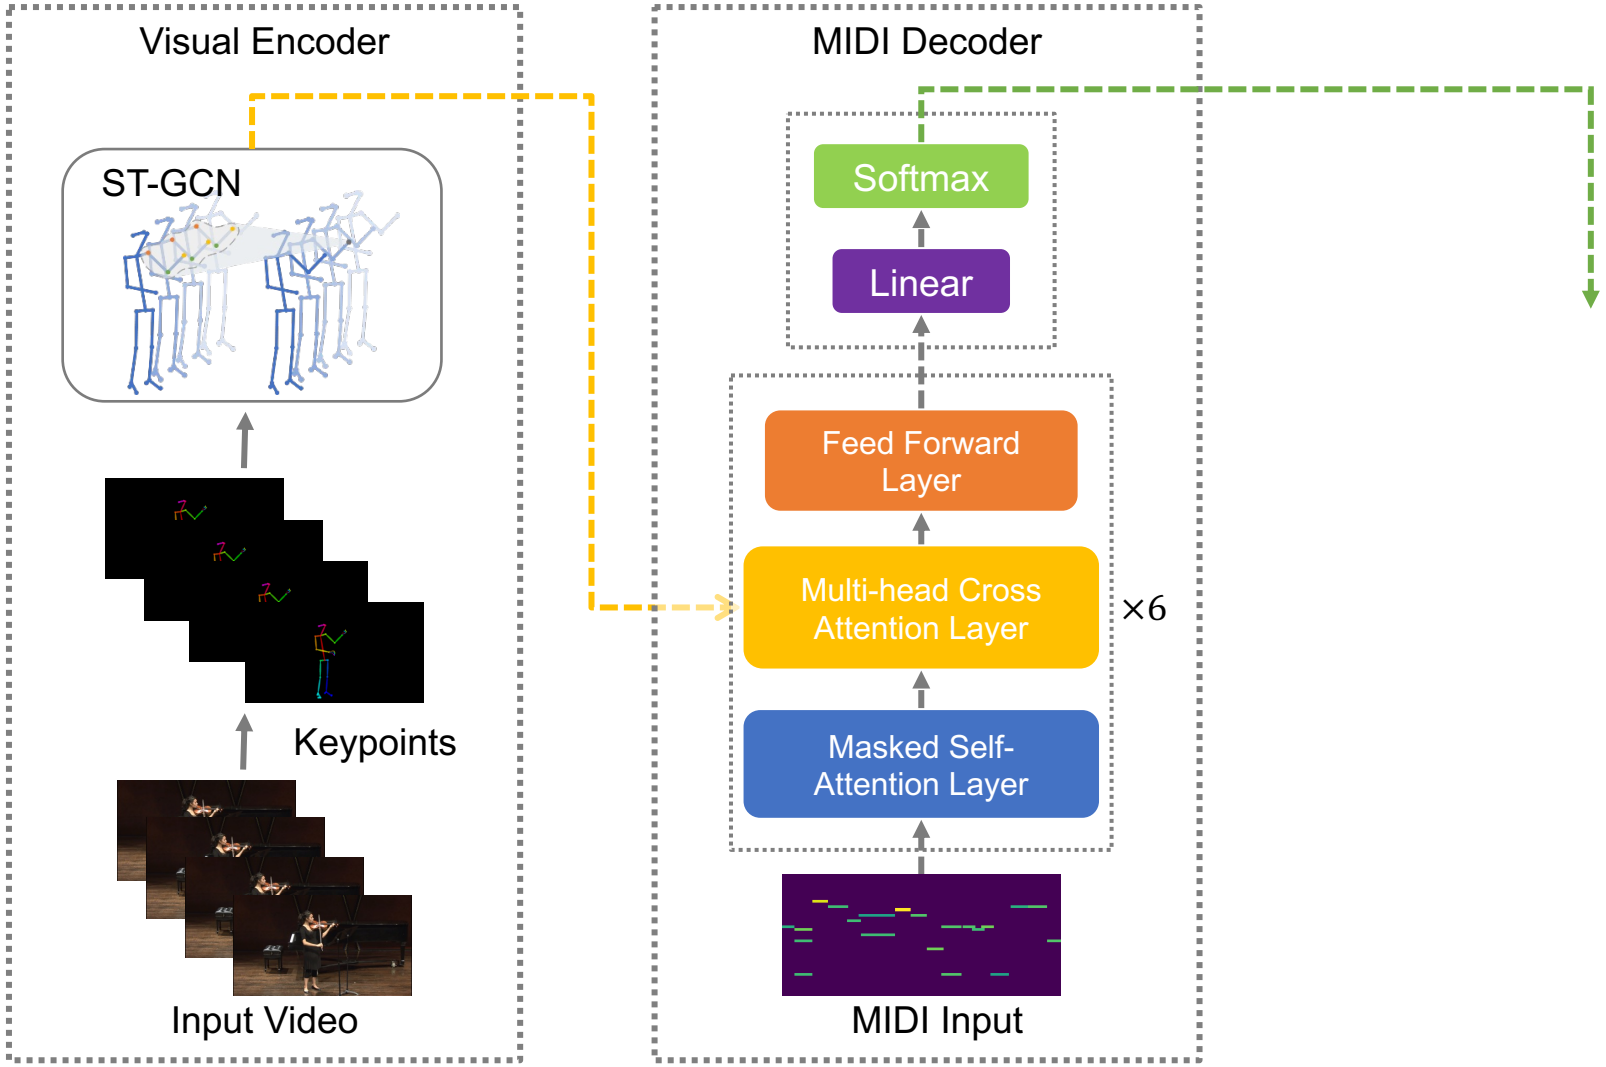

Our method

Finally, we synthesize audio from MIDI.

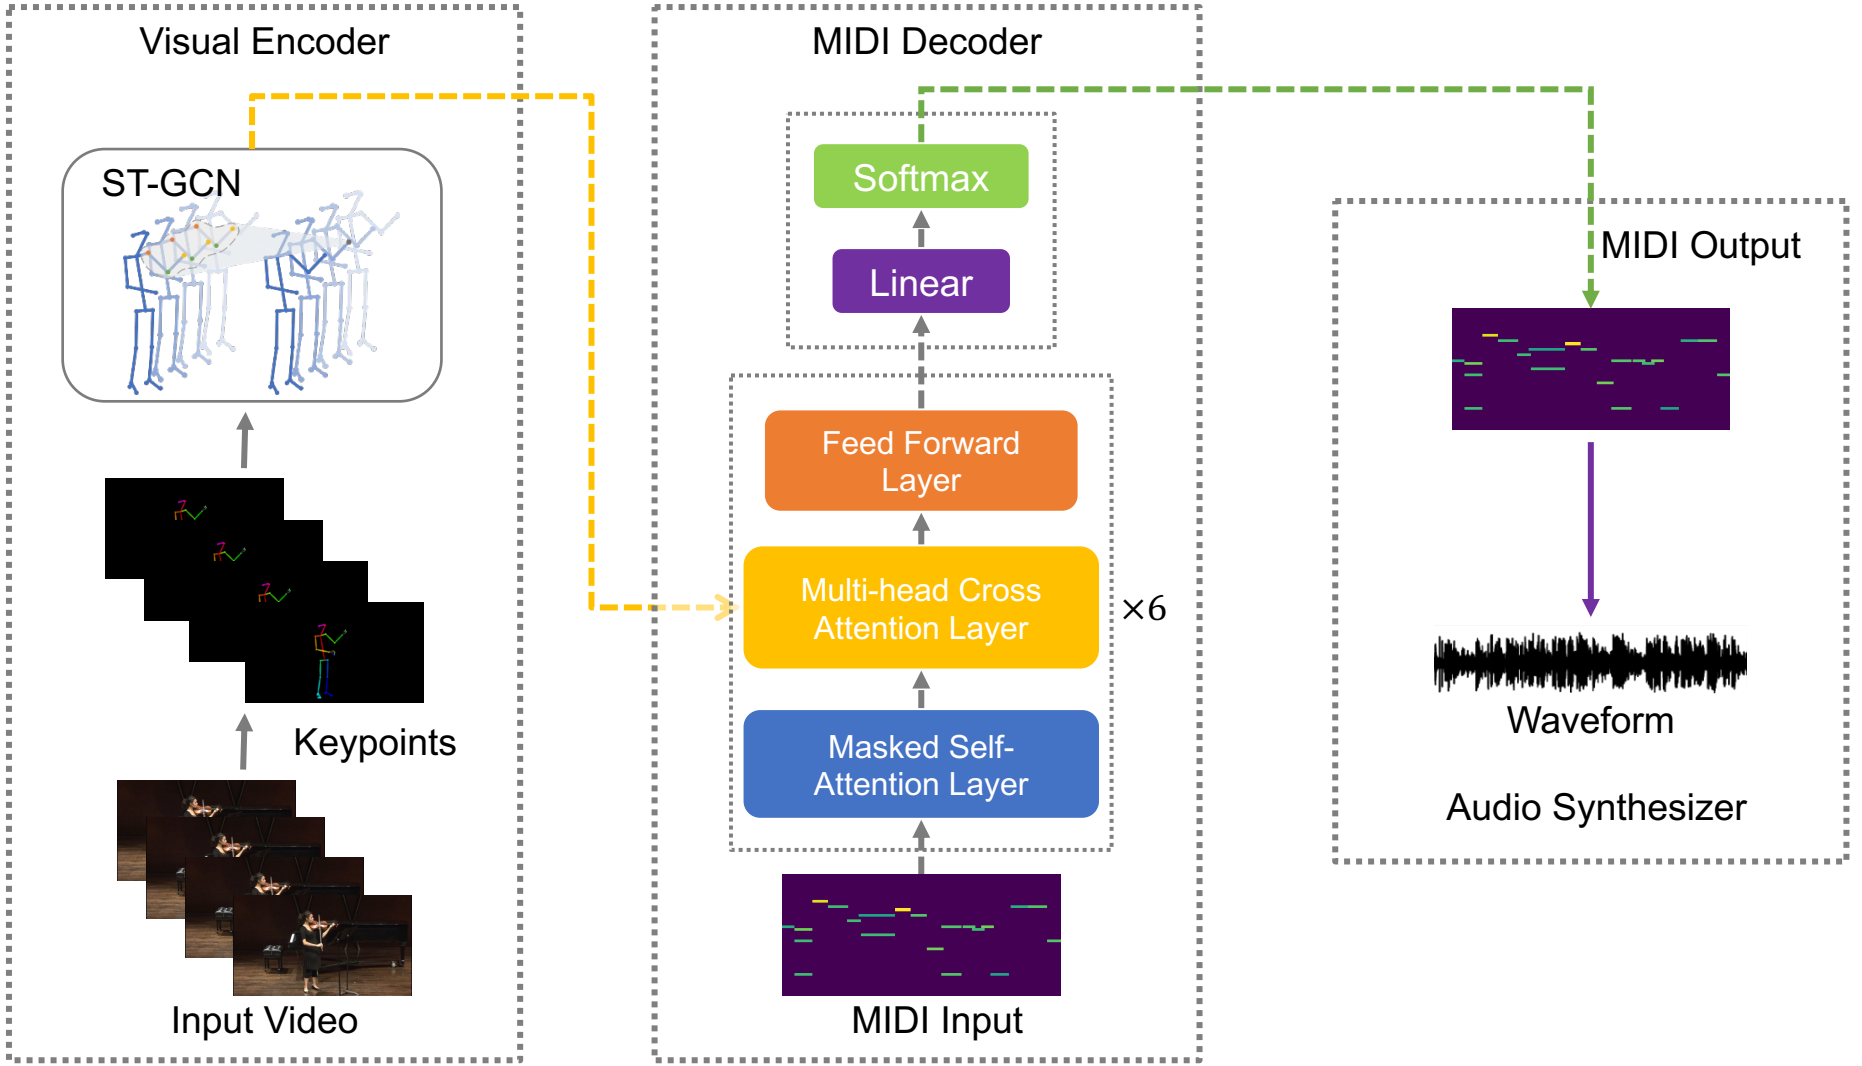

Music generation results

Ukulele

Piano

Guitar

Style editing results

Bass

Original prediction

Style editing

A major

F major

G major

Tuba

Original prediction

Style editing

A major

F major

G major

An impact sound of physical object interactions is critical to perception

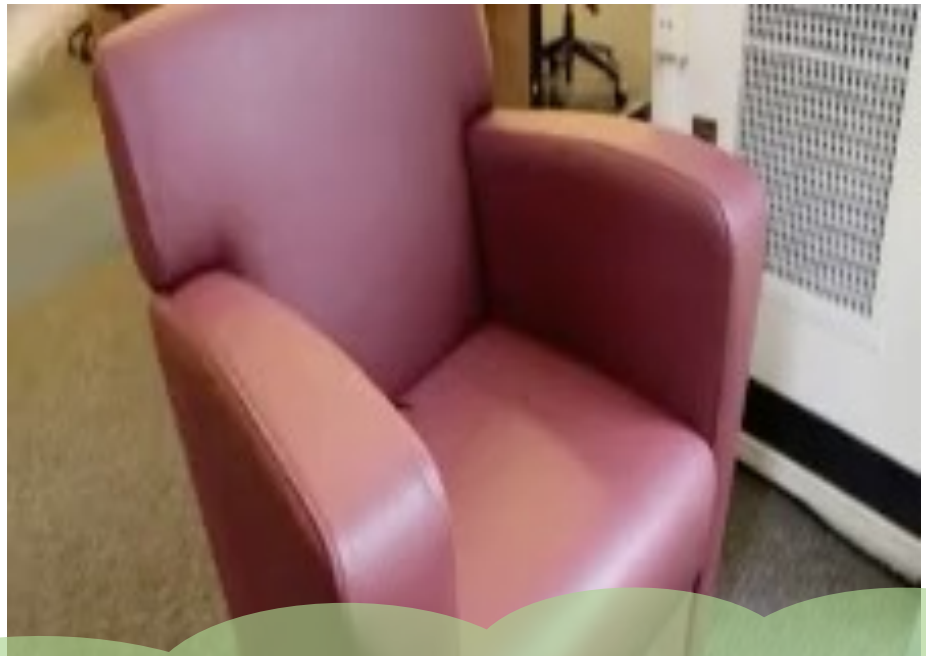

Could we generate the impact sound from vision?

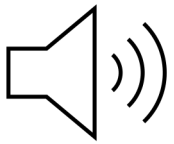

Our Framework

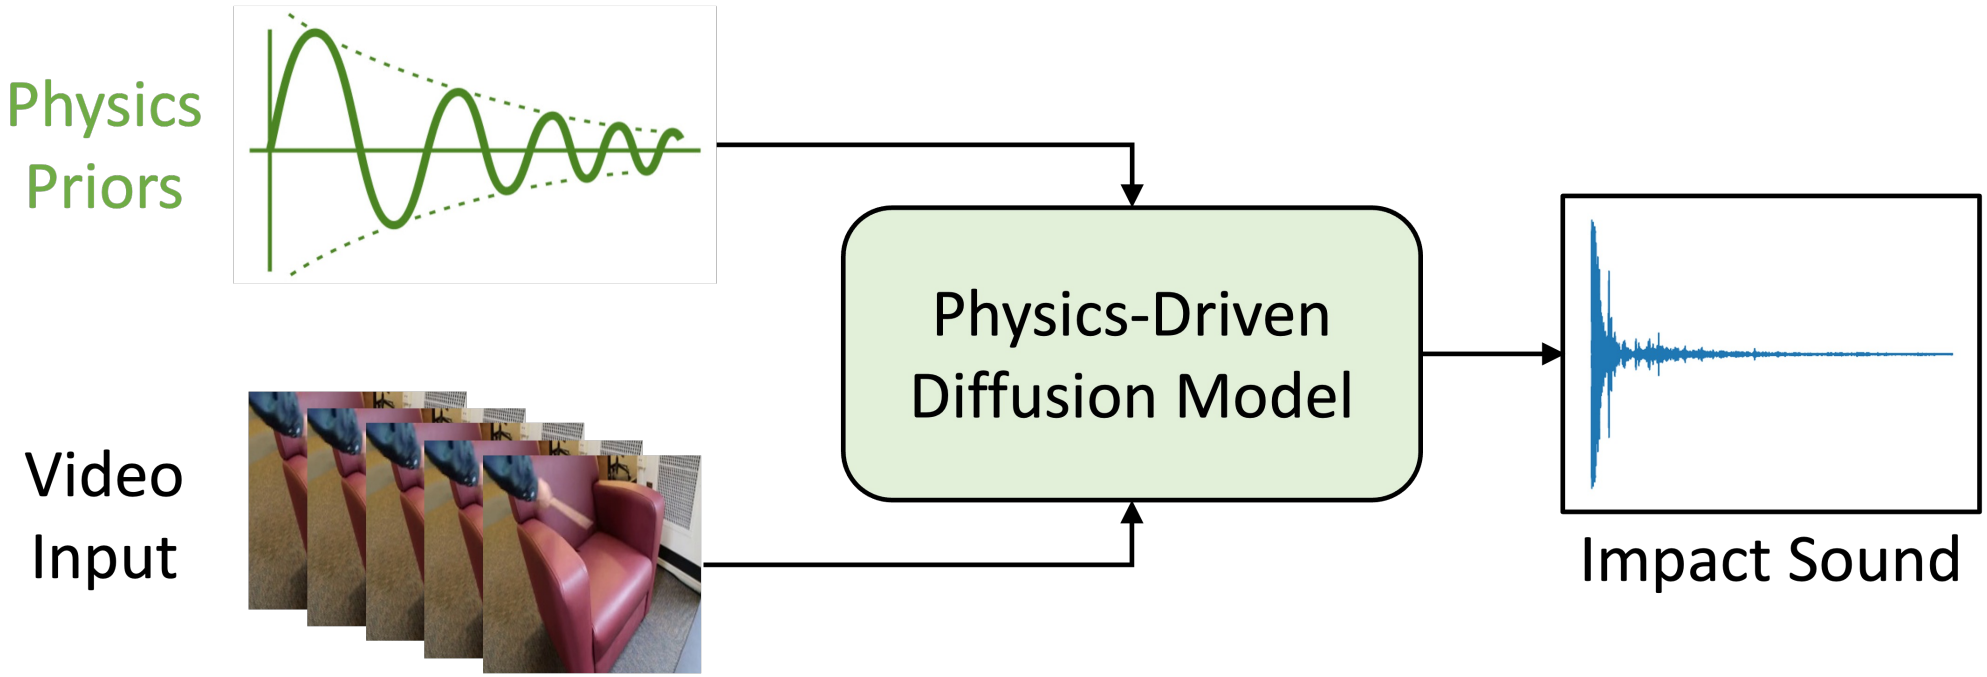

Our model is applicable to
a variety of videos and materials

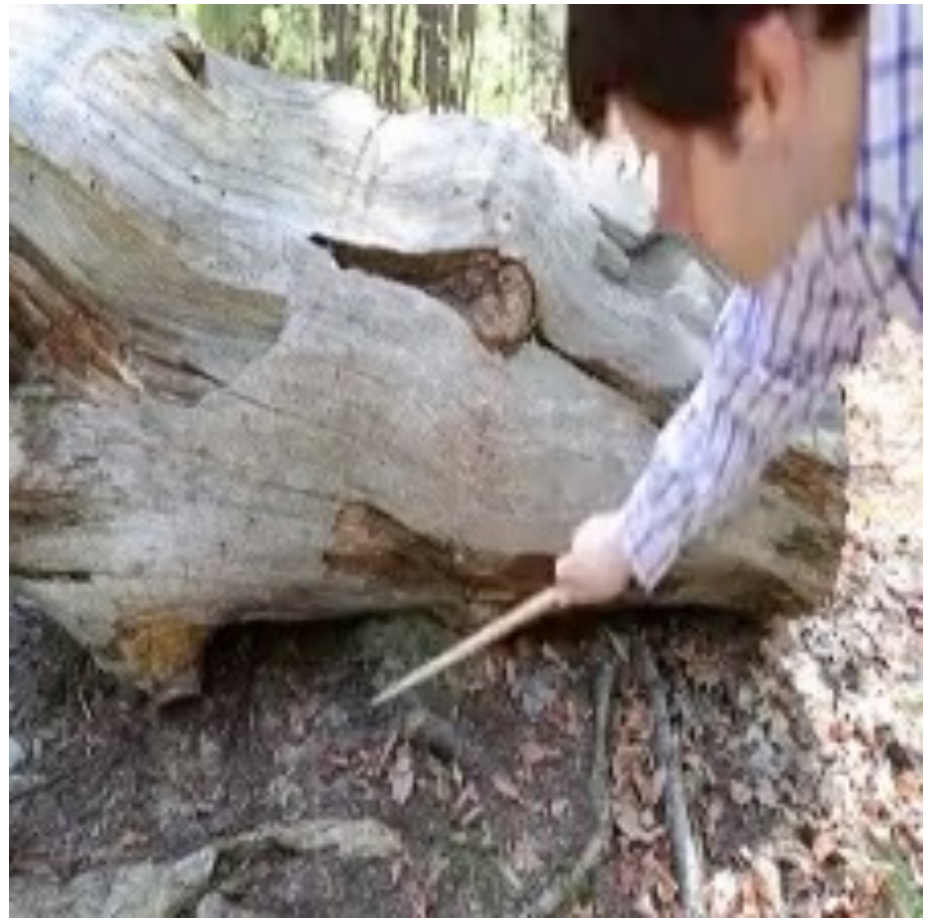

The transparency of **physics priors** allows us to perform interesting **sound editing**

Physics Priors of **Glass** + Video Input

Original Result

Physics Priors of **Glass** + Video Input

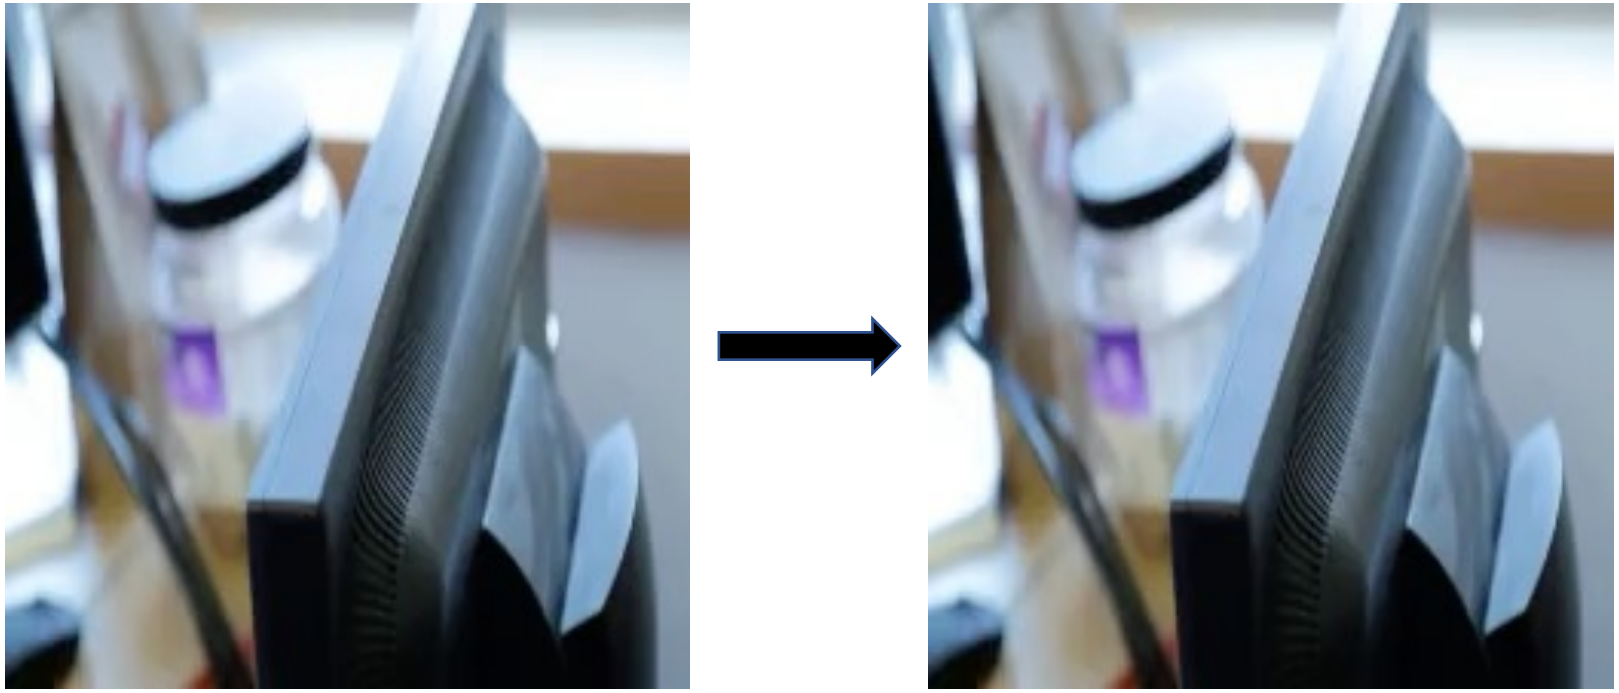

Original Result

Transformed Result

Physics Priors of **Cloth** + Video Input

Original Result

Physics Priors of **Cloth** + Video Input

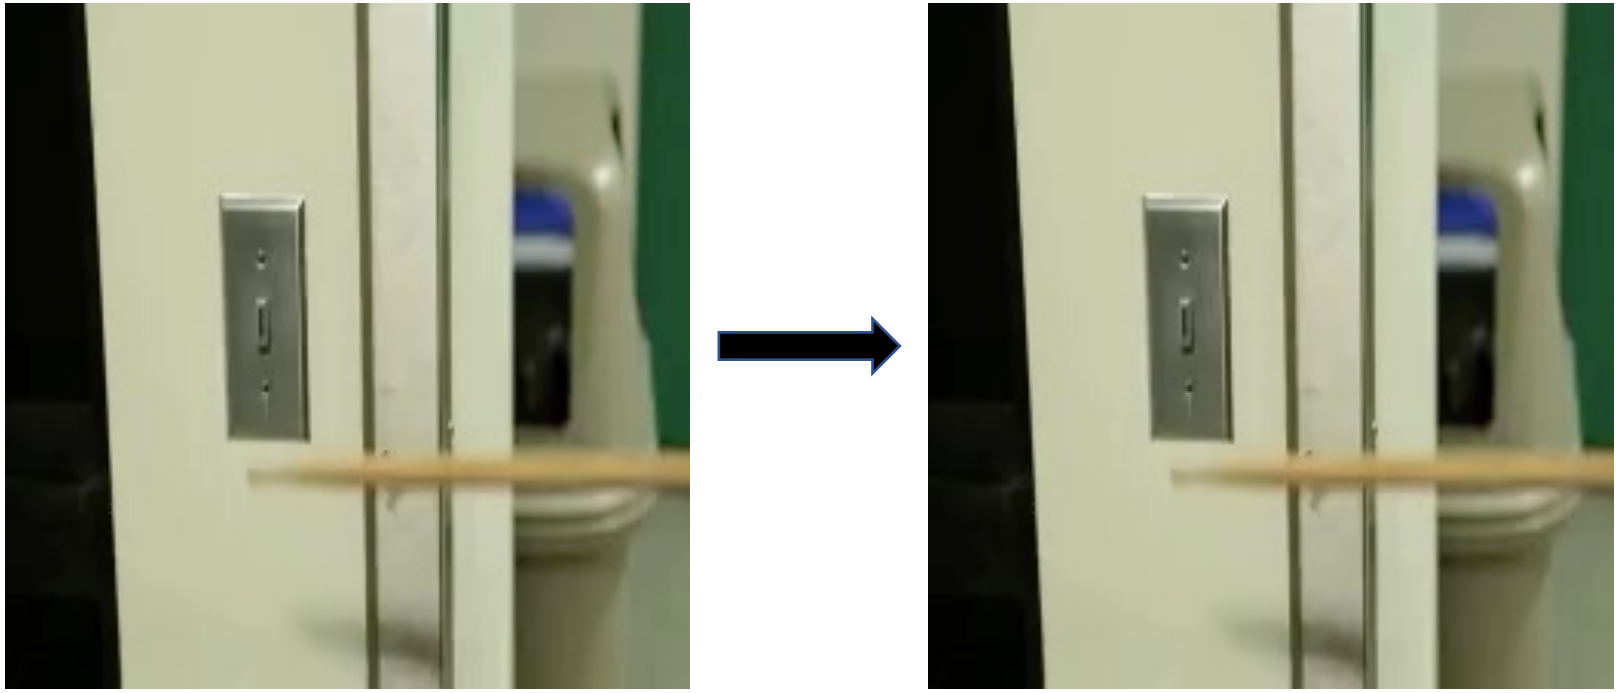

Original Result

Transformed Result

Given an input video…

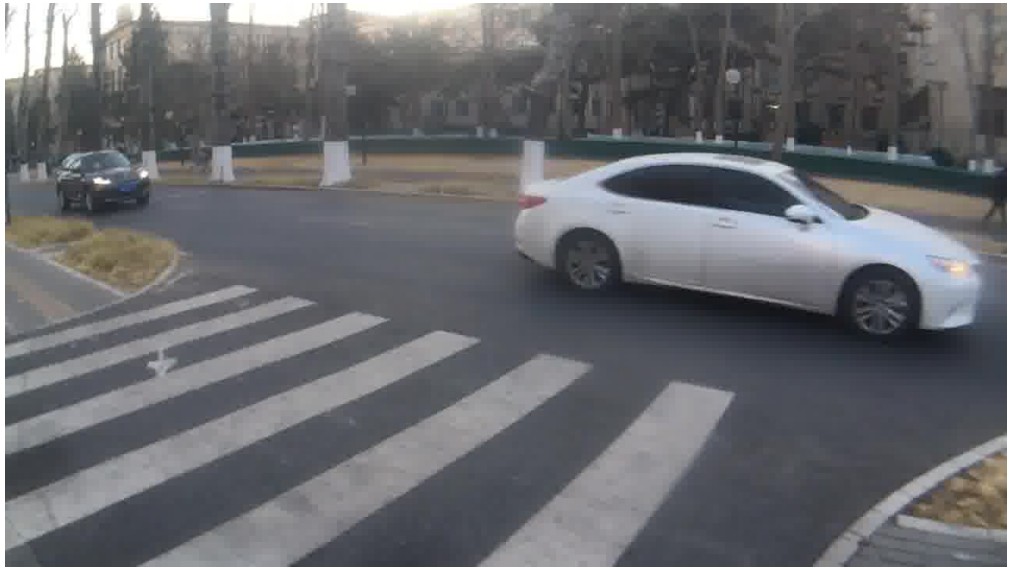

Self-supervised moving vehicle tracking with stereo sound. Gan et.at. ICCV 19

Given an input video…

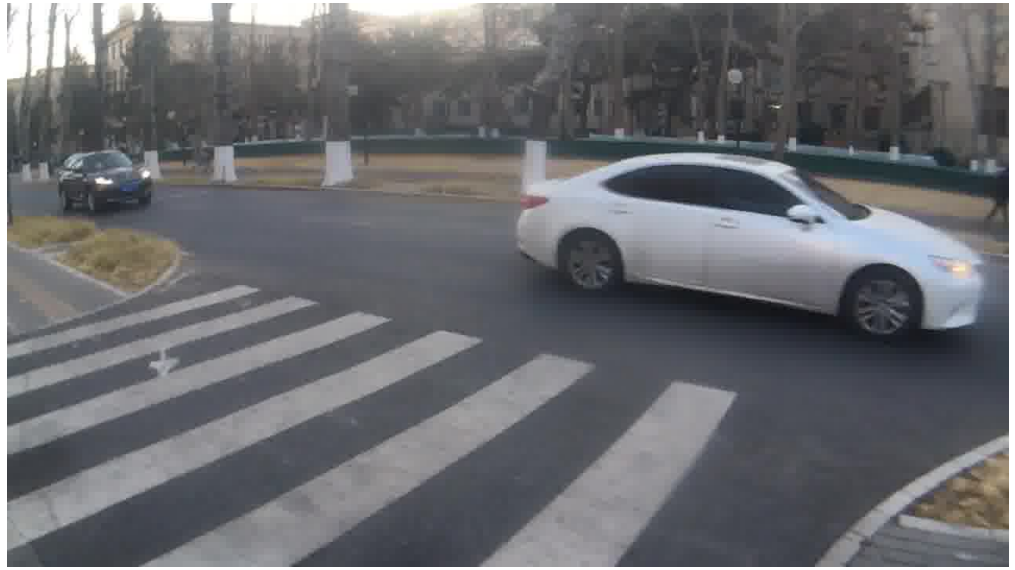

Visual detection network can track the vehicles using visual input.

Visual tracking

What if given a piece of input stereo sound only?

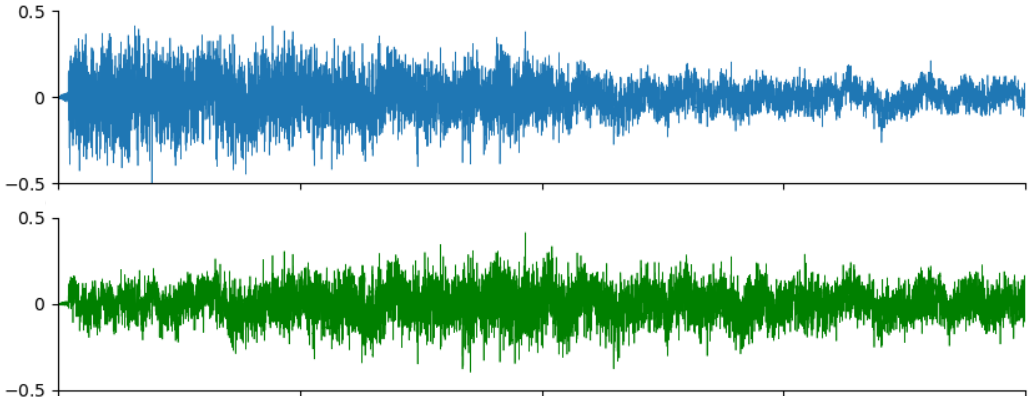

What if given a piece of input stereo sound only?

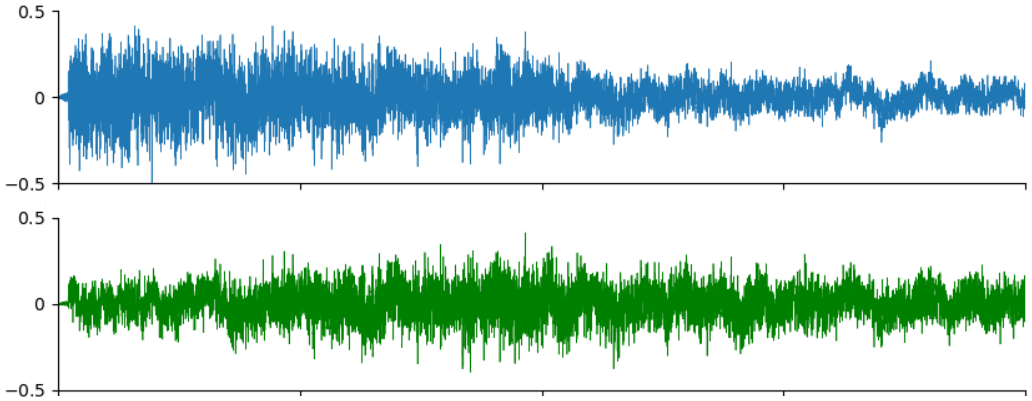

Where are the vehicles?

Applications

- Tracking under poor lighting scene
- Tracking under visual occlusion scene
- Energy-efficiency, privacy-preserving

Our Methods

- Teacher-student alignment

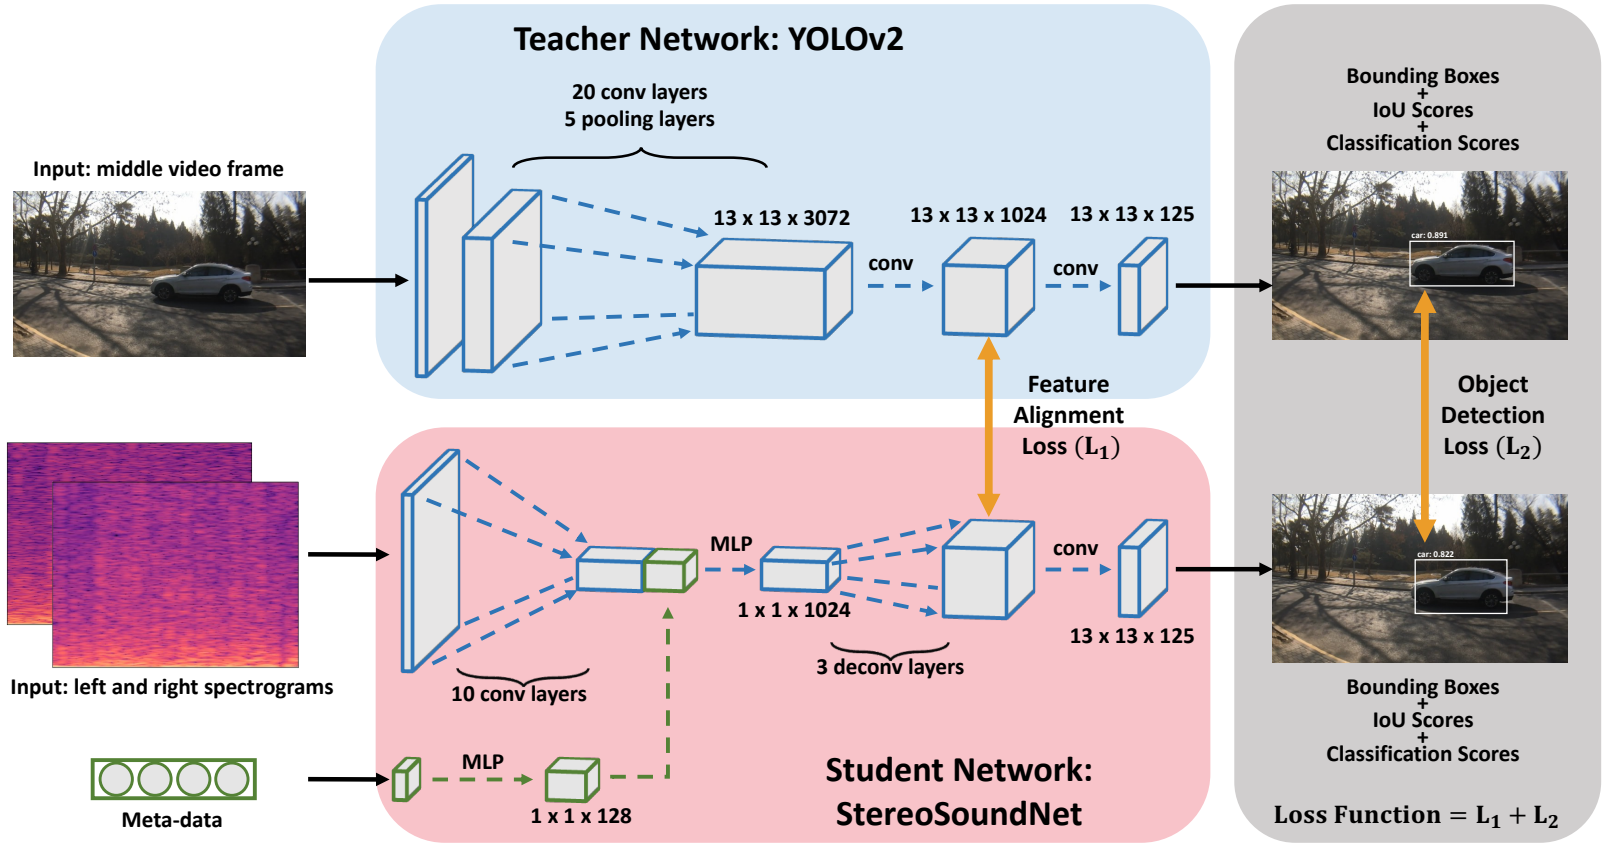

Our Methods

- We first get **pseudo localization labels** and **visual features** from a pre-trained YOLO

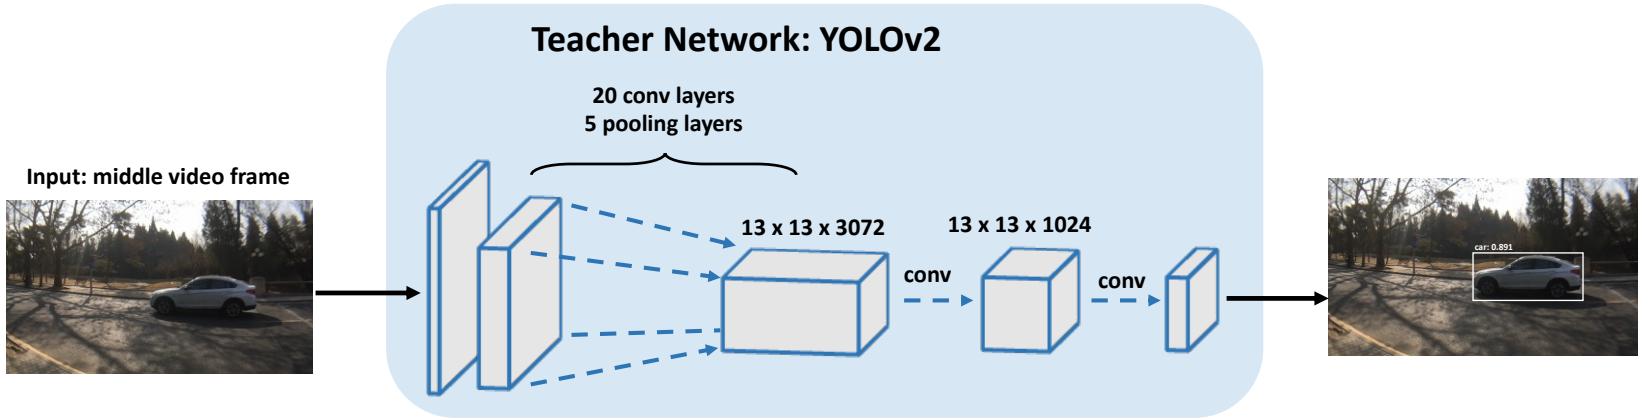

Our Methods

- We then train a sound branch using
  - Feature alignment
  - Object alignment

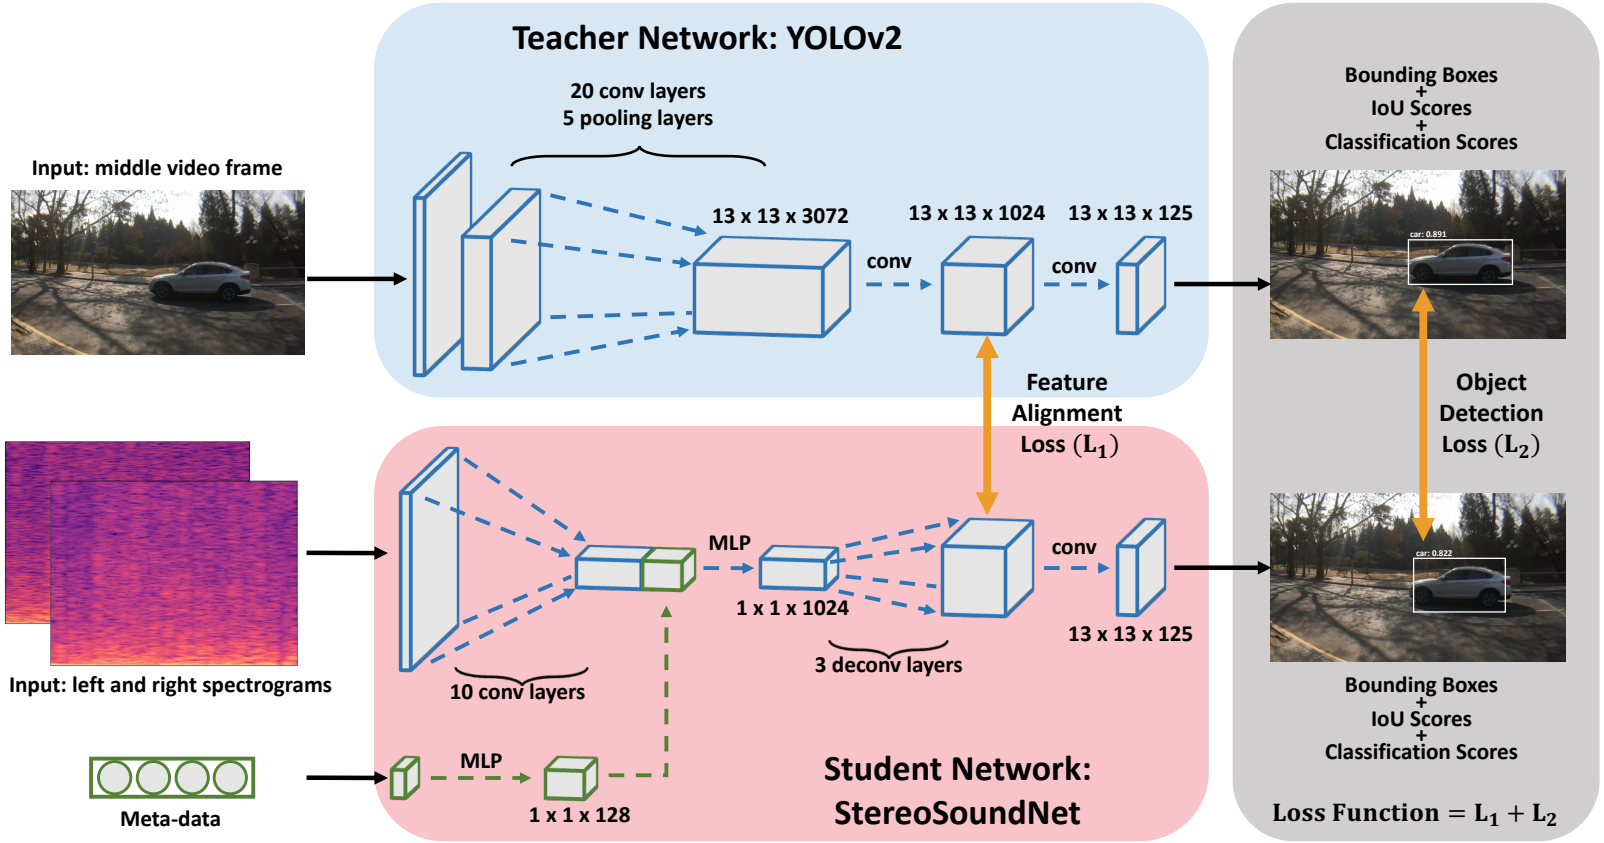

Datasets

- We have collected a dataset on diverse scenes

Results

- Visual tracking fails in many situations:
  - Occlusion
  - Backlighting
  - Reflection on the windows
  - Night scene

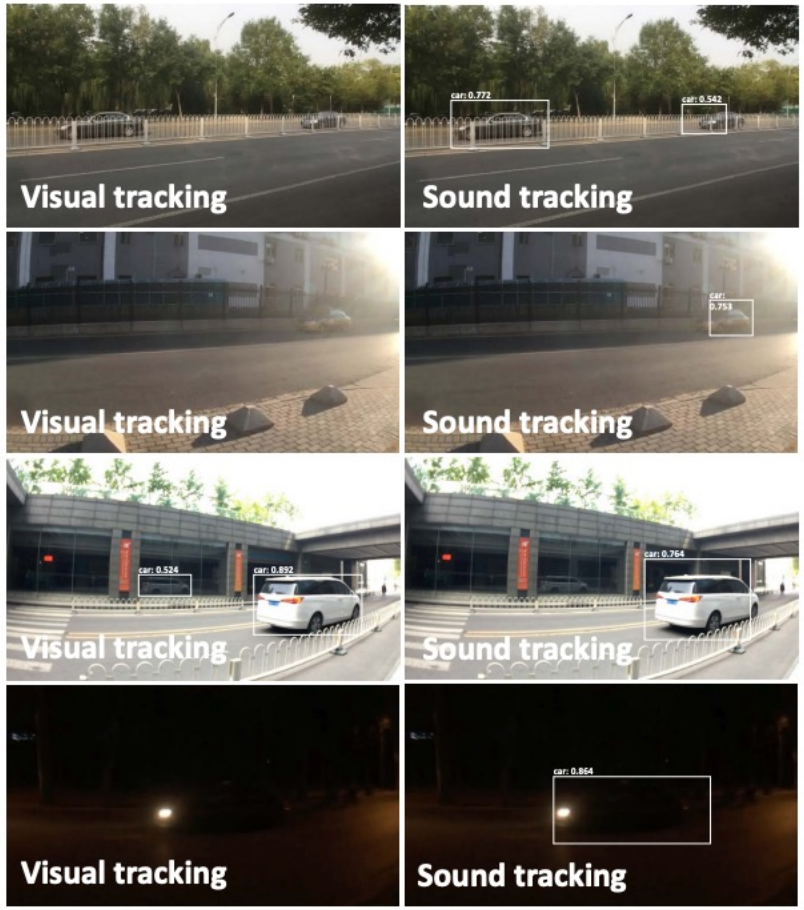

Results

Sound tracking

Tracking Under Poor Lighting

TheeDWorld: 3D Virtual World

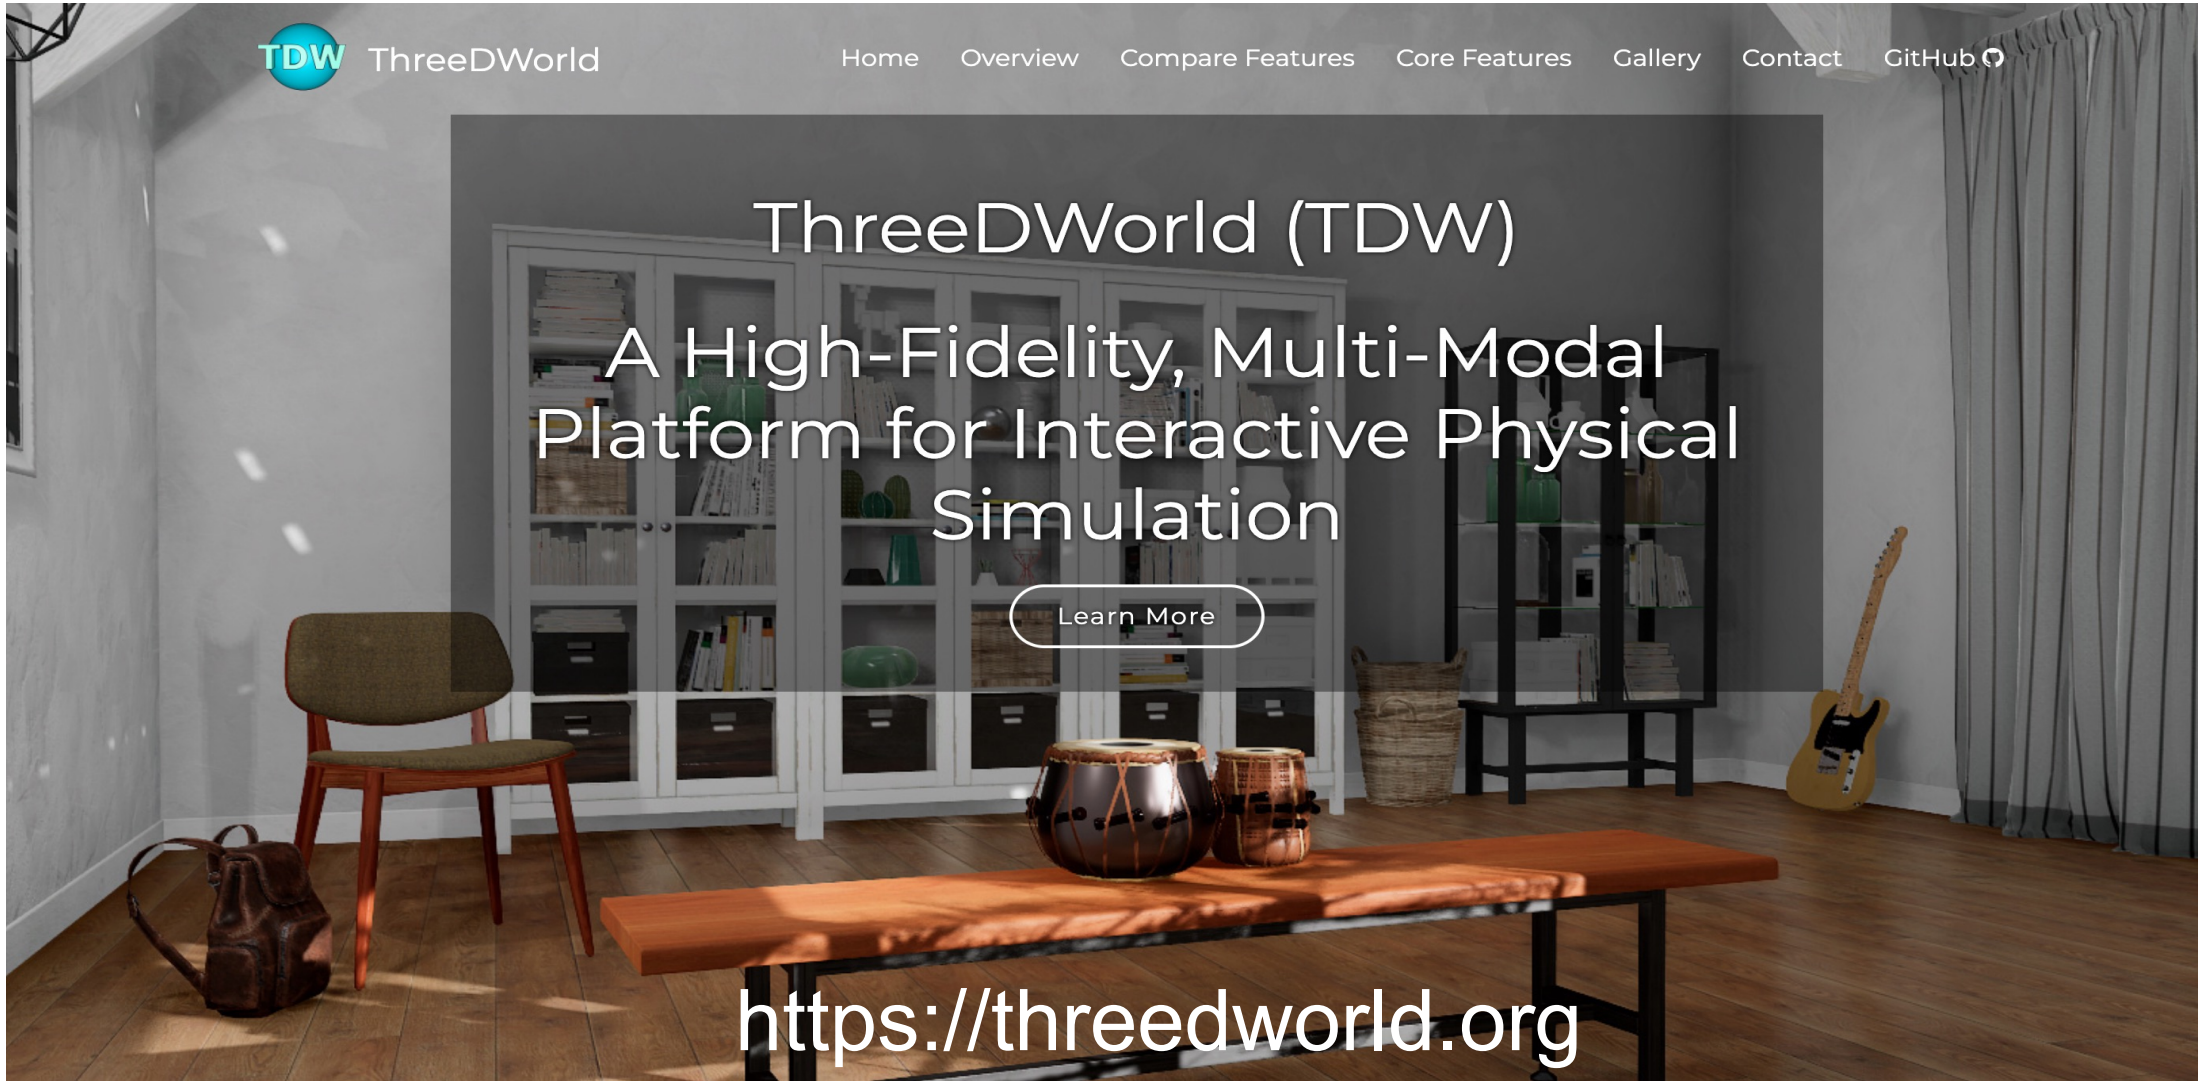

Gan, Jeremy, et al, Threedworld: A platform for interactive multi-modal physical simulation. NeurIPS, 21

Vision: Photo-Realistic Rendering

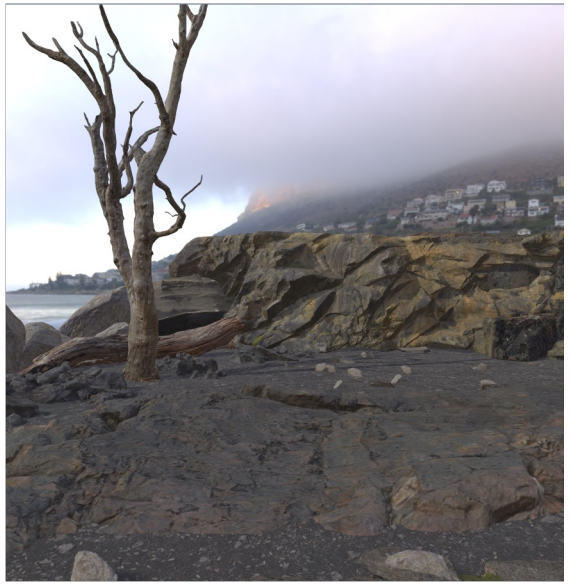

Audio: Physics-Triggered Sound

Physics Simulation

Rigid-body

Soft-body

Agent Interaction

Robot

Humanoid Avatar

Object Interaction in VR

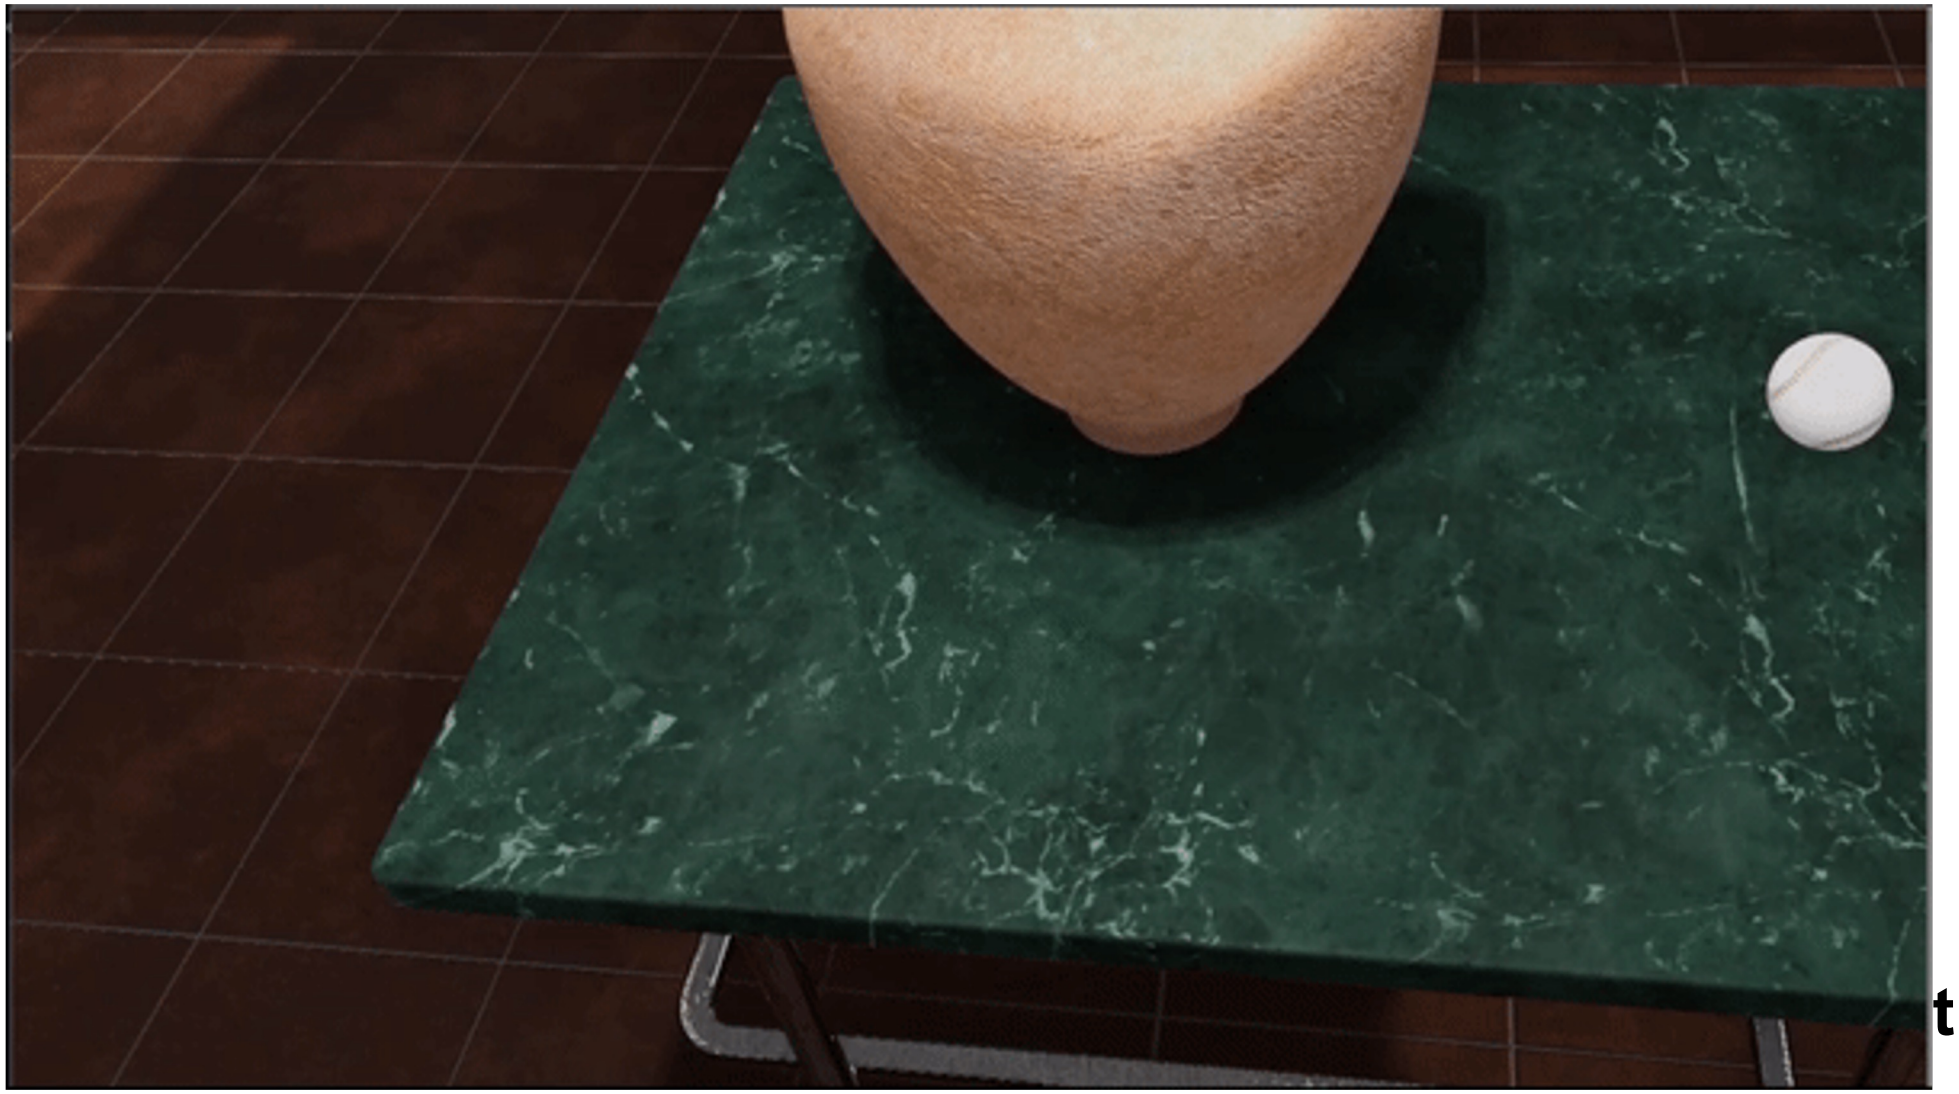

t

Task: Visual Reasoning

Q: What color is the fire hydrant?

A: Yellow

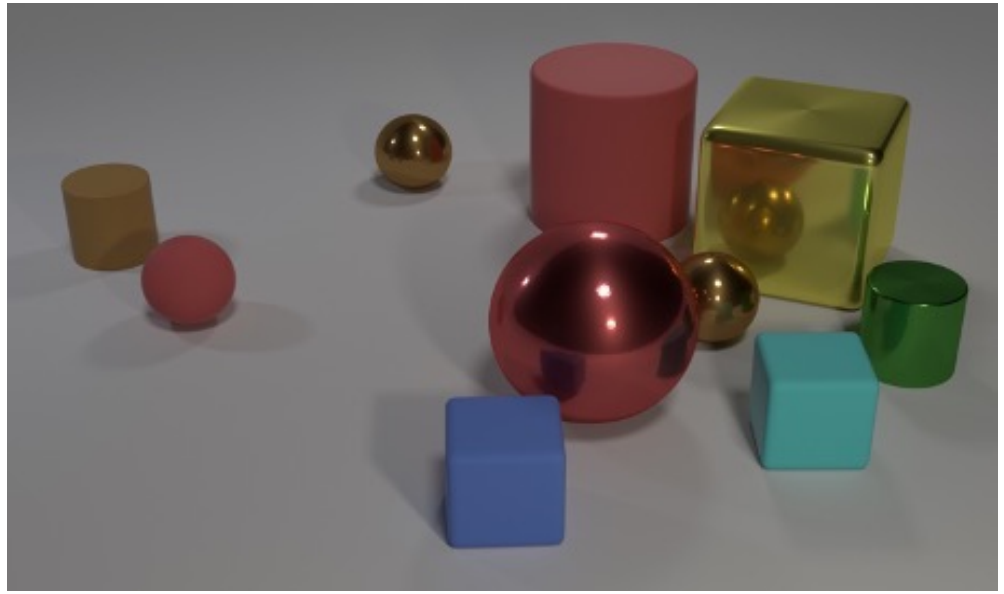

Q: Are there an equal number of large things and metal spheres?

A: Yes

Agrawal et al. VQA, 2015. Johnson et al. CLEVR 2017.

Prior Work: End-to-End Visual Reasoning

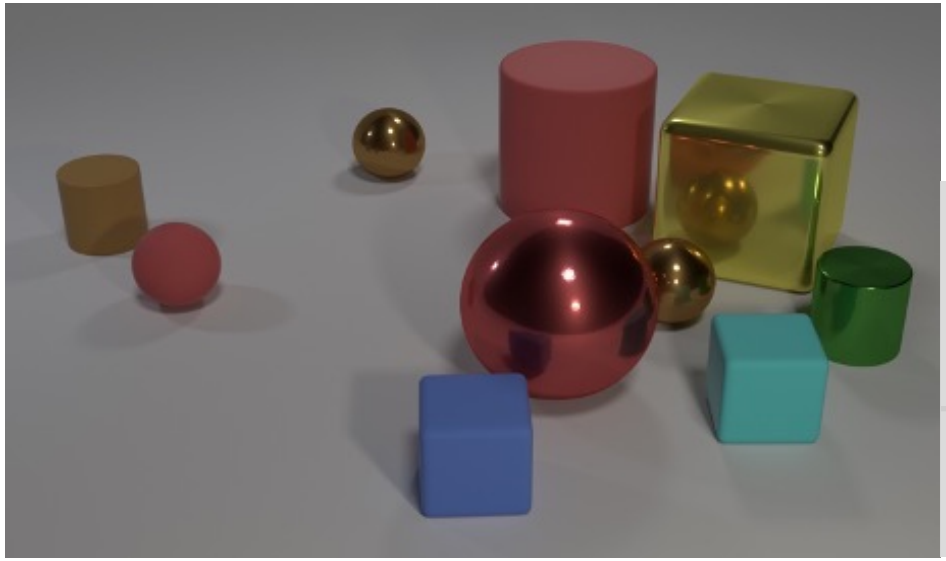

Visual Question Answering

Q: ***Are there an equal number of large things and metal spheres?***

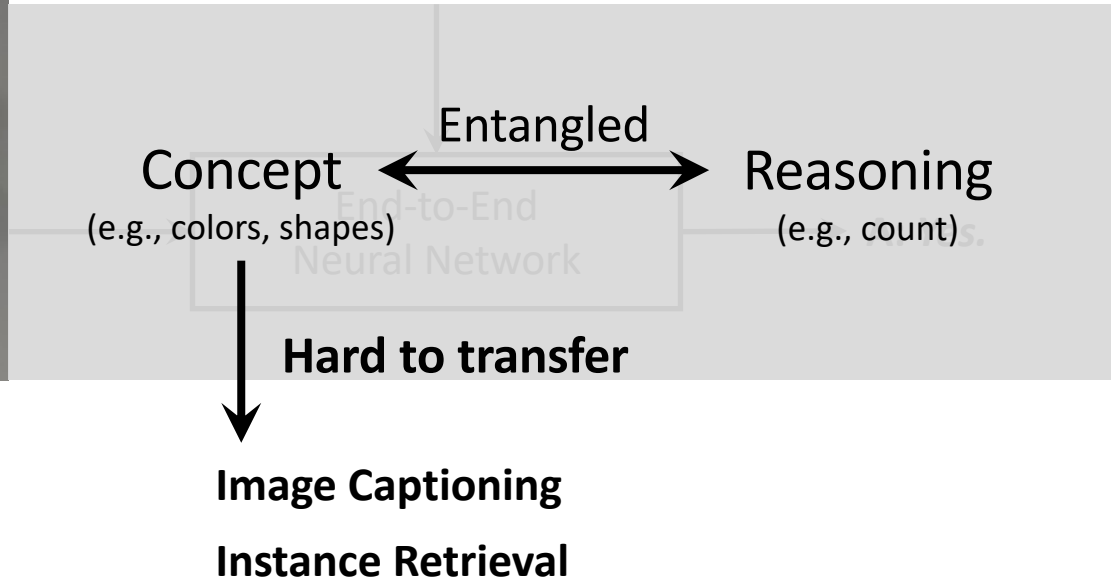

Agrawal et al. VQA, 2015. Johnson et al. CLEVR 2017.
Johnson et al. CLEVR 2017 , Andreas et al. NMN, 2016. Johnson et al. IEP, 2017. Perez et al. FiLM, 2018. Hudson & Manning. MAC, 2018. Hu et al. Stack-NMN, 2018. Mascharka et al. TbD, 2018.

How Do Human Reason From A Visual Scene?

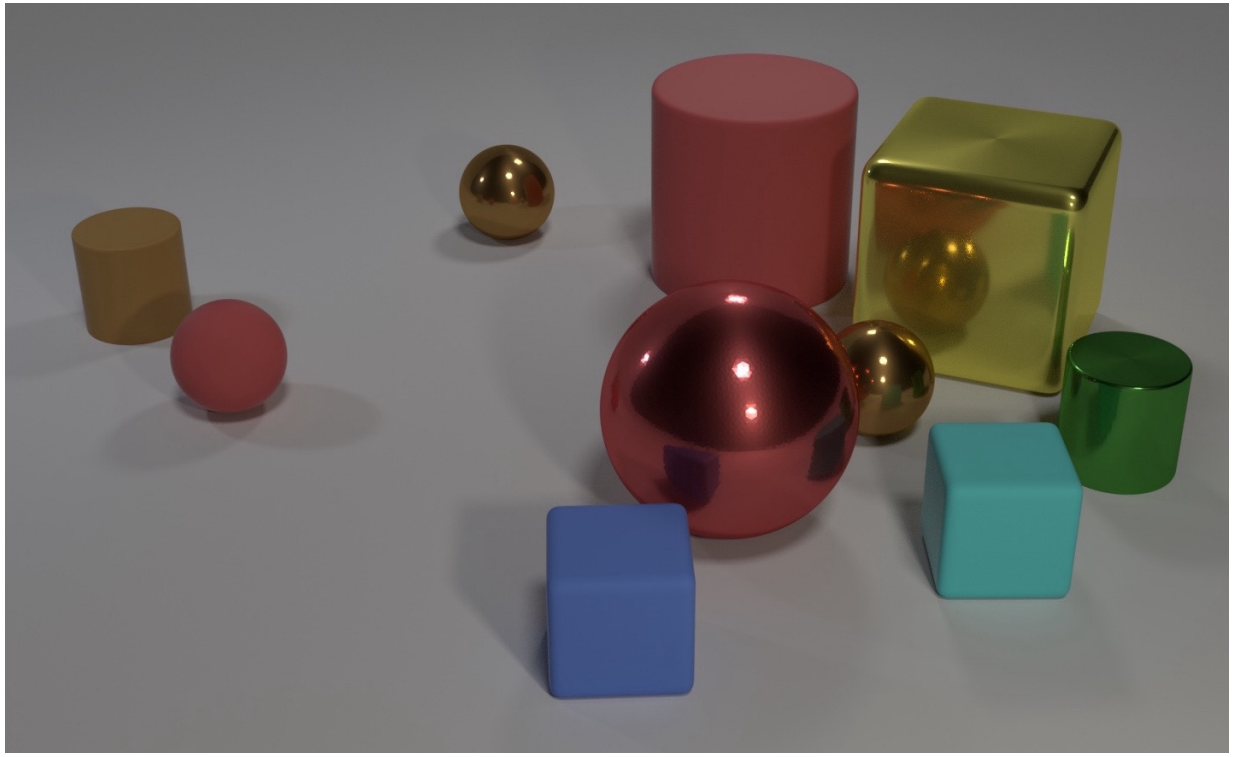

Question: *Are there an equal number of large things and metal spheres?*

How Do Human Reason From A Visual Scene?

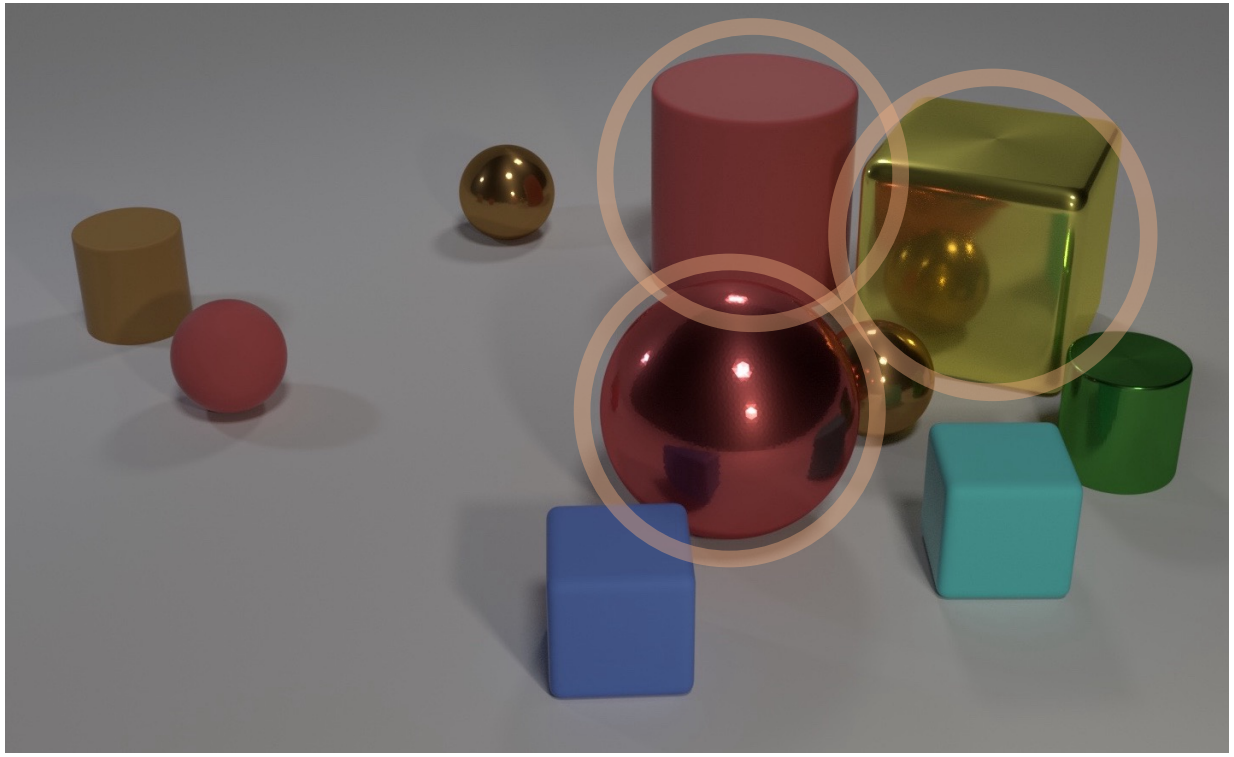

Question: *Are there an equal number of large things and metal spheres?*

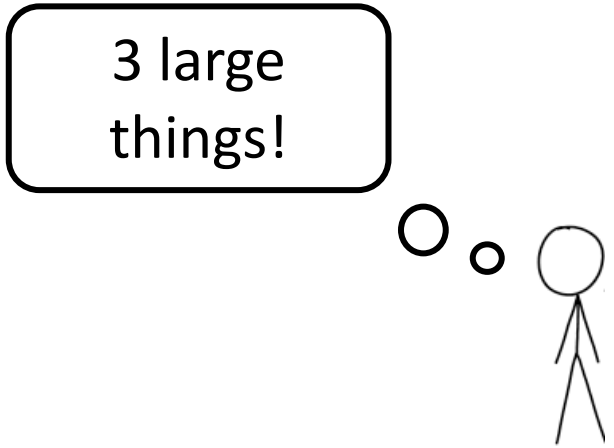

How Do Human Reason From A Visual Scene?

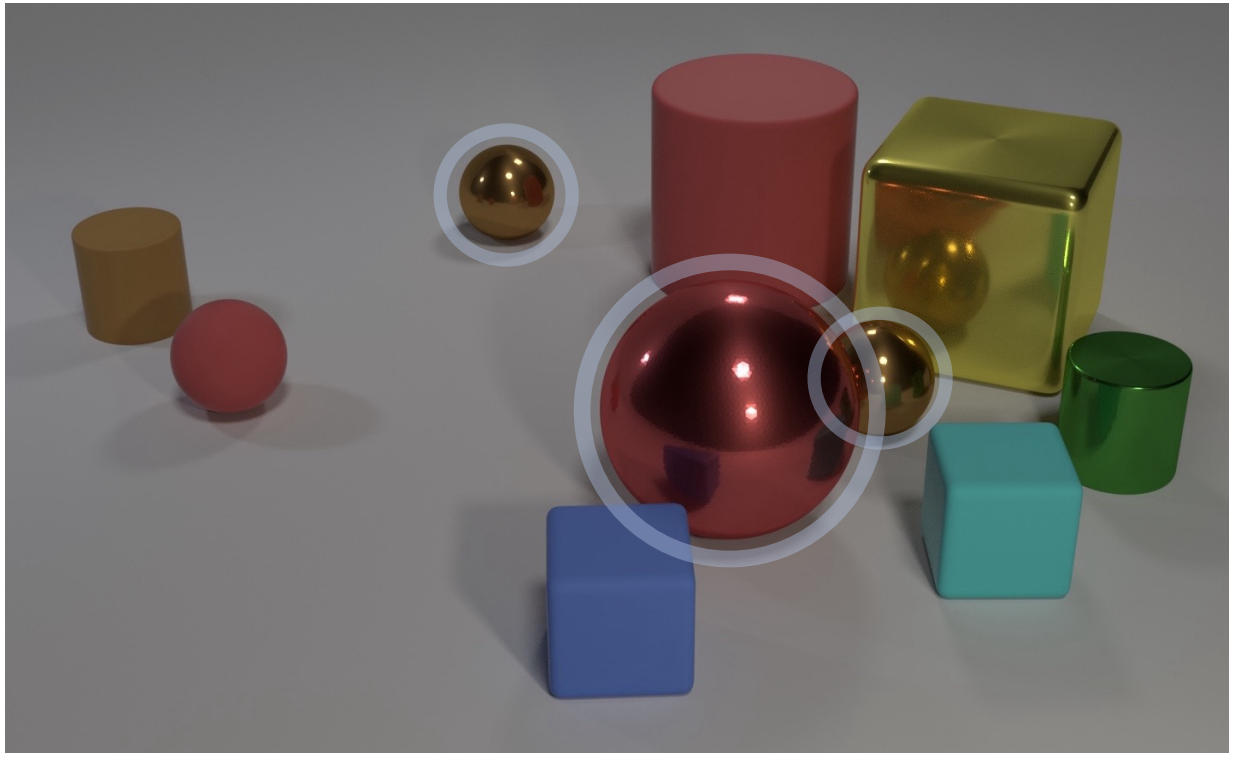

Question: *Are there an equal number of large things and metal spheres?*

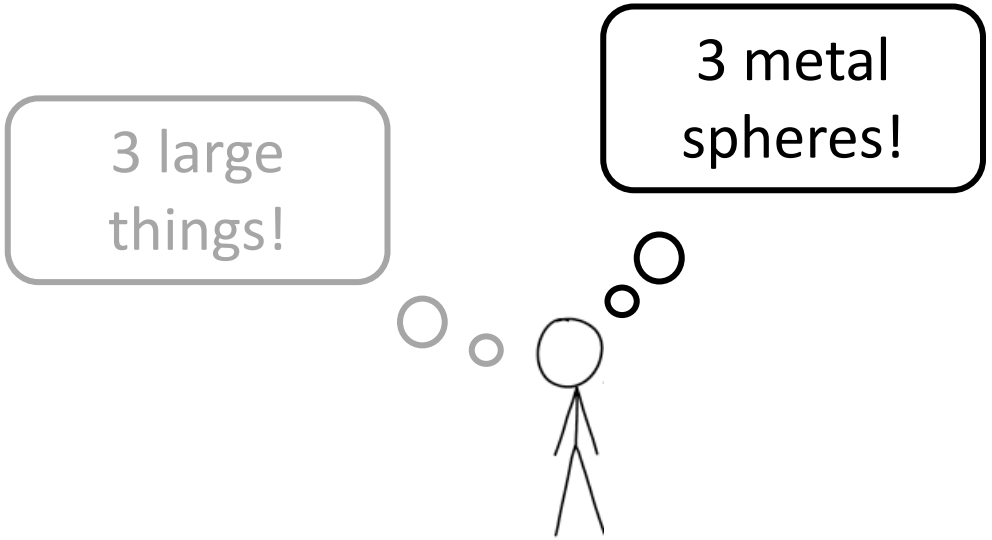

How do human reason from a visual scene?

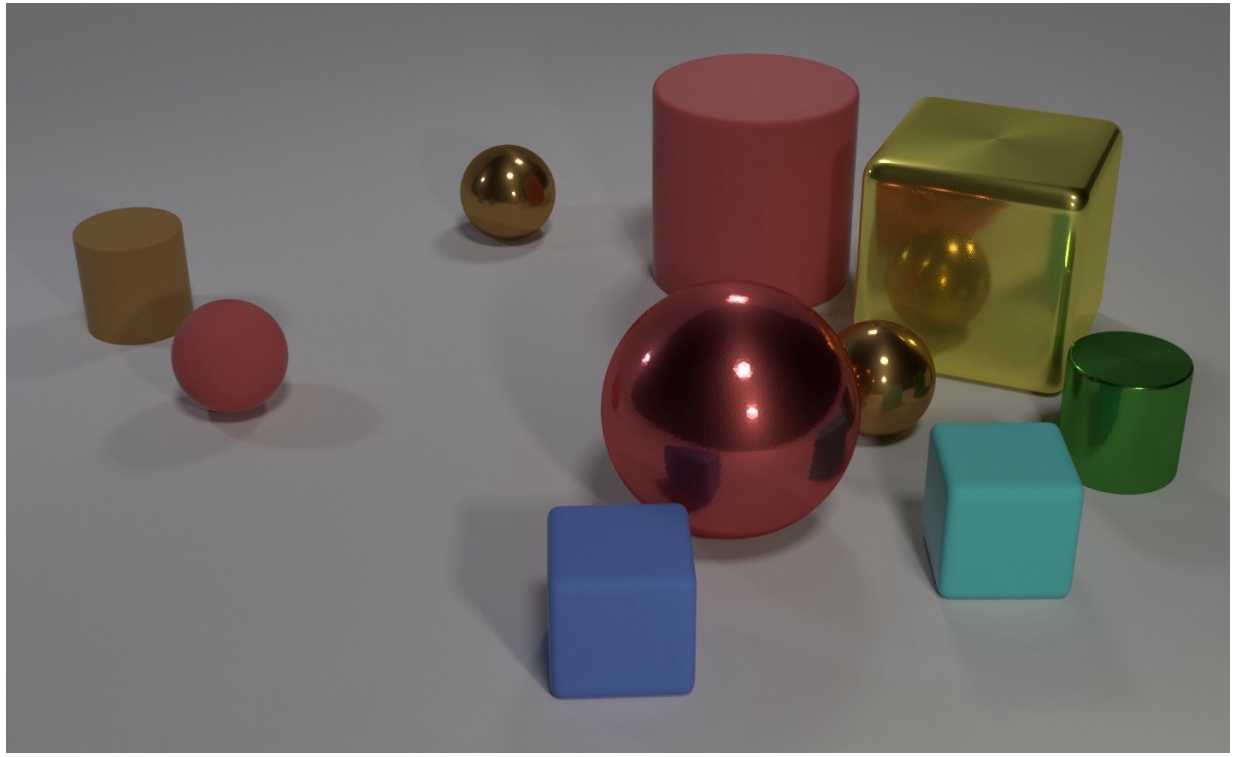

Question: *Are there an equal number of large things and metal spheres?*

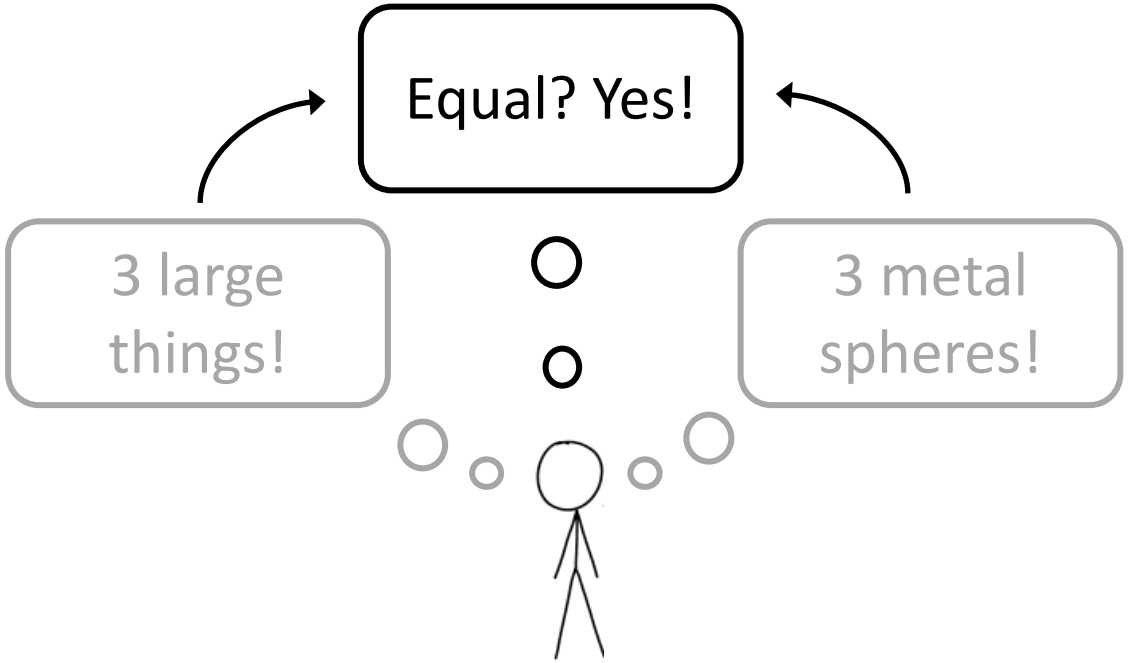

How Do Human Reason From A Visual Scene?

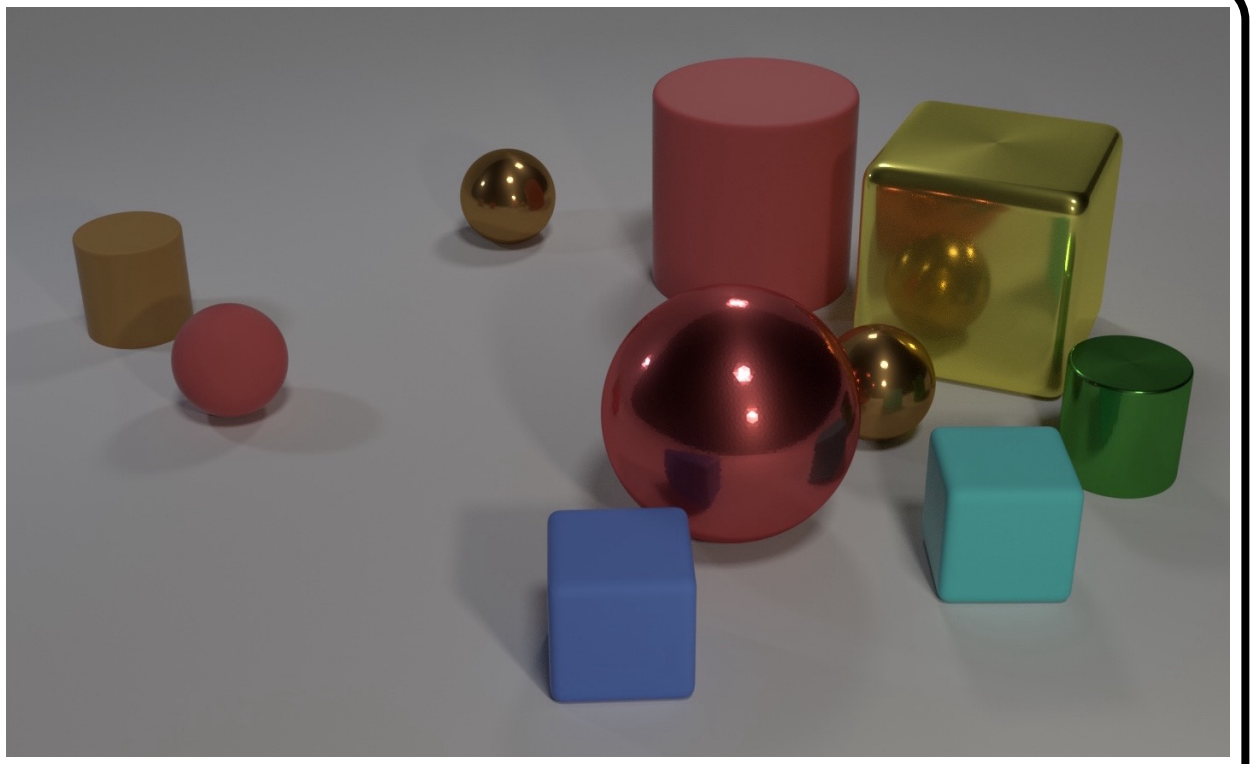

Visual Perception

Question Understanding

Question: *Are there an equal number of large things and metal spheres?*

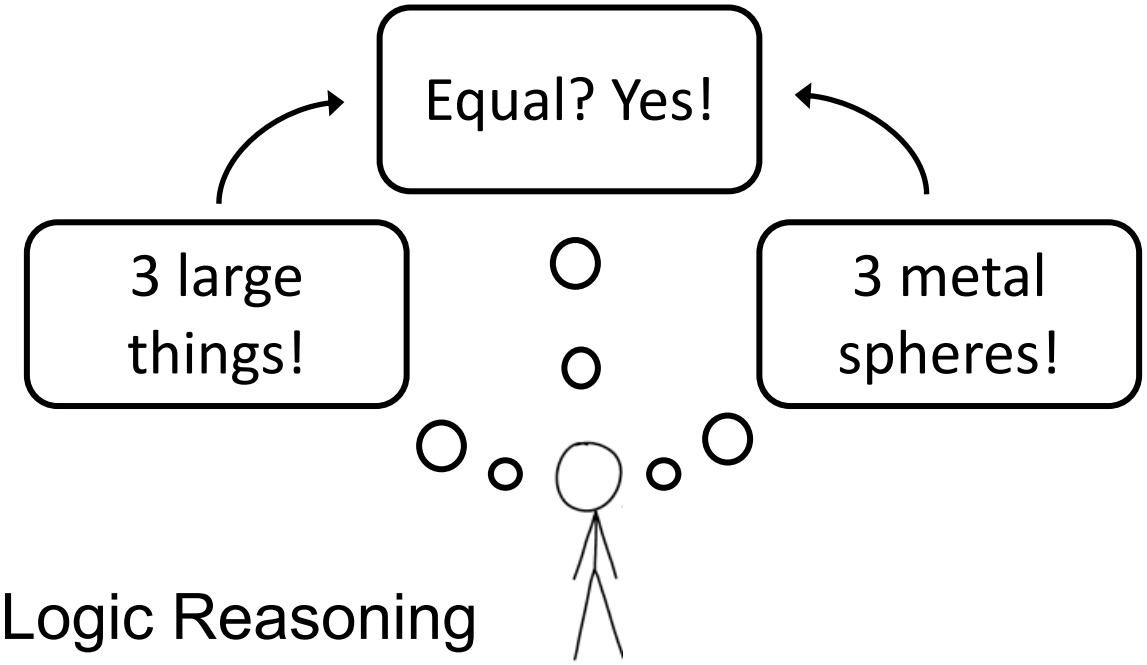

Logic Reasoning

Incorporate Concepts and Symbolic Programs

Neural networks parse images in symbolic concepts

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	0.85	-1.10	0.35
2	Small	Cube	Metal	Purple	1.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

Neural networks parse questions into symbolic functional program.

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind
3. filter_cube
4. filter_large
5. count

Executing programs on the symbolic space.

VQS: Linking Segmentations to Questions and Answers for VQA. **Gan** et al. ICCV'17
Neural-Symbolic VQA: Disentangling Reasoning from Vision and Language Understanding. Yi, Wu, **Gan**, et al. NeurIPS'18

Neural Scene Parsing

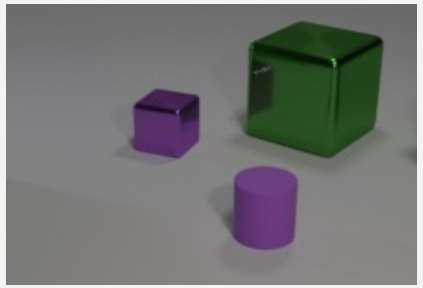

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? →

LSTM Encoder

LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
2	Small	Cylinder	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

5. count

Answer: 1

Neural Scene Parsing

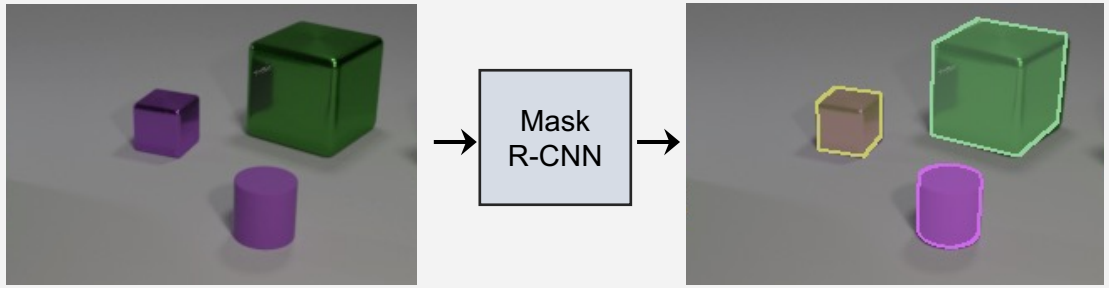

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large?

→ LSTM Encoder → LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
2	Small	Cylinder	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

5. count

Answer: 1

Neural Scene Parsing

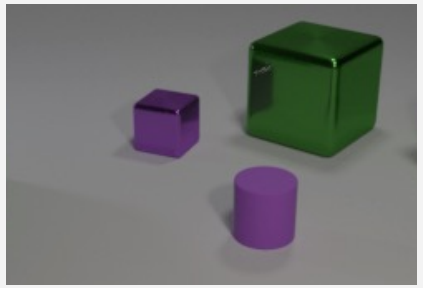

→ Mask R-CNN →

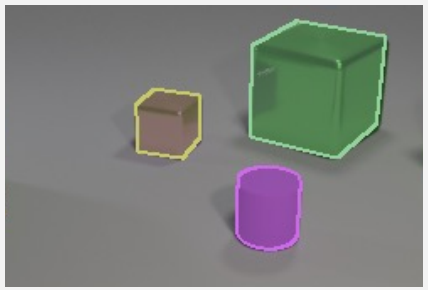

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder →

1. filter_shape(scene, cylinder)
2. relate(behind)
3. filter_shape(scene, cube)
4. filter_size(scene, large)
5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
2	Small	Cylinder	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

5. count

Answer: 1

Neural Question Parsing

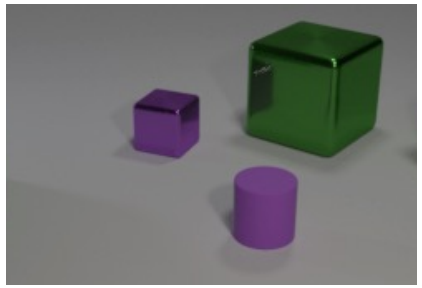

→ Mask R-CNN →

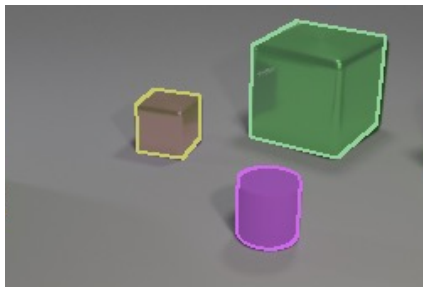

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? →

LSTM Encoder | LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
2	Small	Cylinder	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

5. count

Answer: 1

Symbolic Reasoning

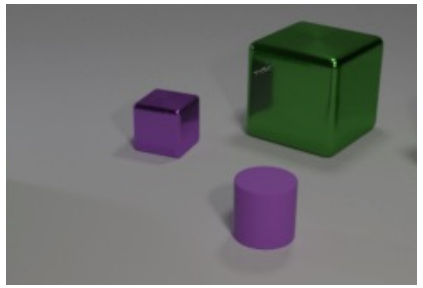

→ Mask R-CNN →

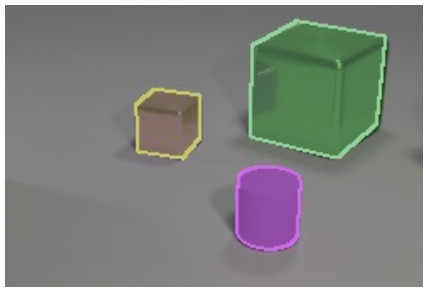

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

III. *Symbolic* Program Execution

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder →

1. filter_shape(scene, cylinder)
2. relate(behind)
3. filter_shape(scene, cube)
4. filter_size(scene, large)
5. count(scene)

Symbolic Reasoning

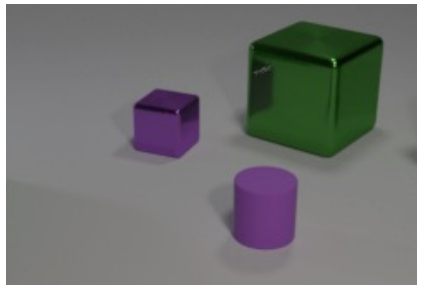

→ Mask R-CNN →

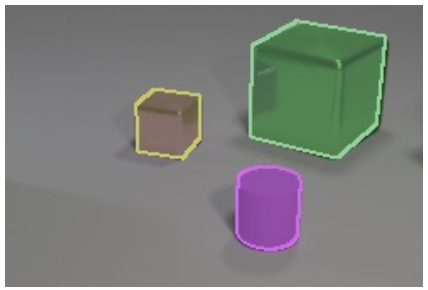

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

↓

III. *Symbolic* Program Execution

1. filter_cylinder

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder →

1. filter_shape(scene, cylinder)
2. relate(behind)
3. filter_shape(scene, cube)
4. filter_size(scene, large)
5. count(scene)

→

Symbolic Reasoning

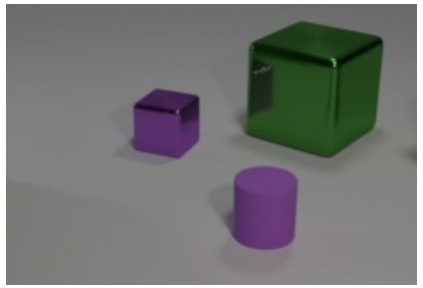

→ Mask R-CNN →

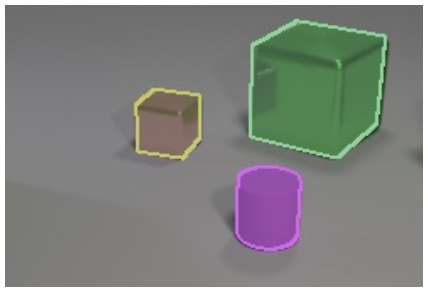

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? →

LSTM Encoder | LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

→

ID	Size	Shape	…
1	Small	Cube	…
3	Large	Cube	…

Symbolic Reasoning

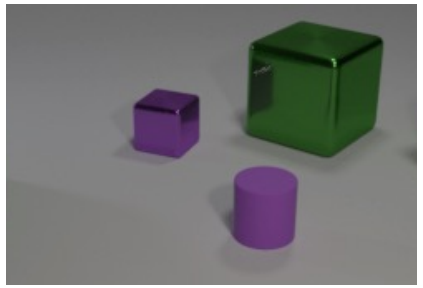

→ Mask R-CNN →

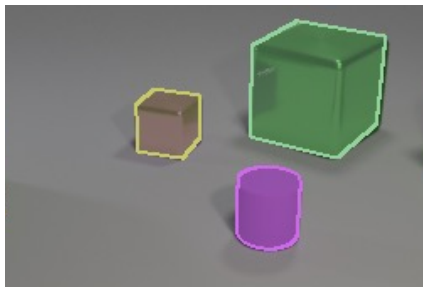

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

III. *Symbolic* Program Execution

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
3	Large	Cube	…

3. filter_cube

Symbolic Reasoning

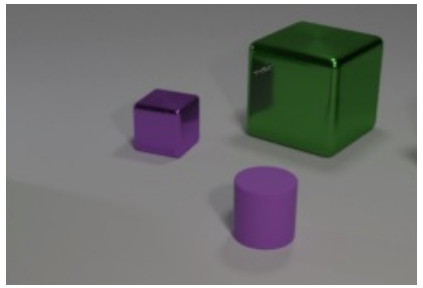

→ Mask R-CNN →

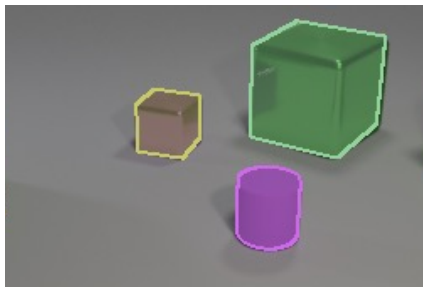

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder →

1. filter_shape(scene, cylinder)
2. relate(behind)
3. filter_shape(scene, cube)
4. filter_size(scene, large)
5. count(scene)

III. *Symbolic* Program Execution

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

Symbolic Reasoning

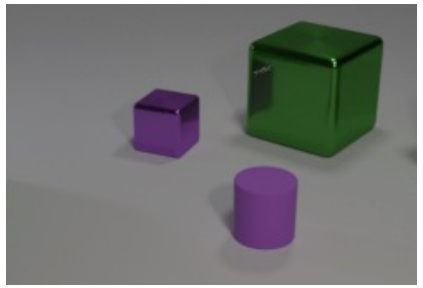

→ Mask R-CNN →

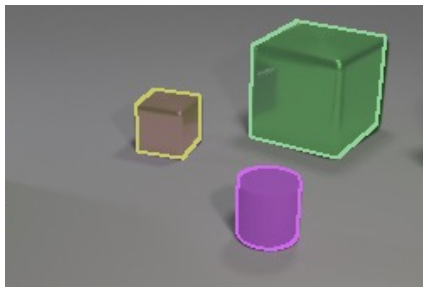

→ CNN →

ID	Size	Shape	Material	Color	x	y	z
1	Small	Cube	Metal	Purple	-0.45	-1.10	0.35
2	Small	Cylinder	Rubber	Purple	0.75	1.31	0.35
3	Large	Cube	Metal	Green	1.58	-1.60	0.70

I. *Neural* Scene Parsing

II. *Neural* Question Parsing

III. *Symbolic* Program Execution

How many cubes that are behind the cylinder are large? → LSTM Encoder → LSTM Decoder

→ 1. filter_shape(scene, cylinder)
→ 2. relate(behind)
→ 3. filter_shape(scene, cube)
→ 4. filter_size(scene, large)
→ 5. count(scene)

1. filter_cylinder
2. relate_behind

ID	Size	Shape	…
1	Small	Cube	…
3	Large	Cube	…

3. filter_cube
4. filter_large

ID	Size	…
3	Large	…

5. count

Answer: 1

Evaluation on CLEVR

Method	Accuracy (%)
Human	92.6
RN	95.5
IEP	96.9
FiLM	97.6
MAC	98.9
NS-VQA (Ours)	**99.8**

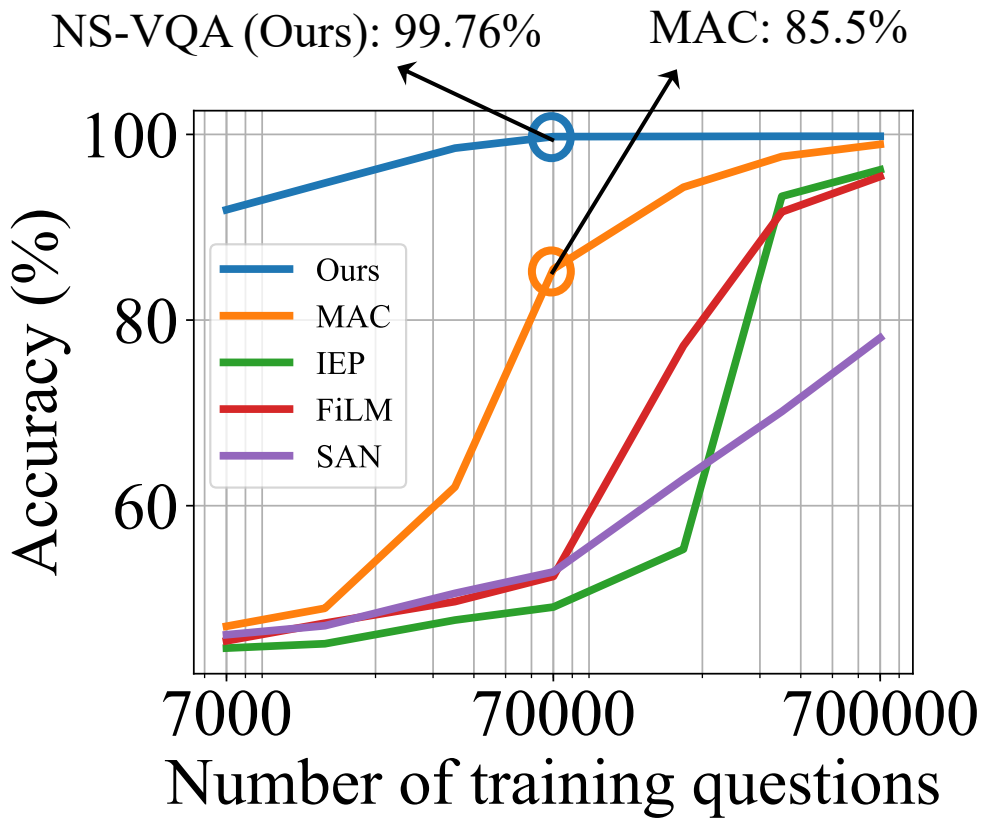

[Johnson et al. ICCV 2017, Santoro et al. NIPS 2017, Perez et al. AAAI 2018, Hudson et al. ICLR 2018, Mascharka et al. CVPR 2018]